

COMPUTER SCIENCE

Database Management System

Query Language

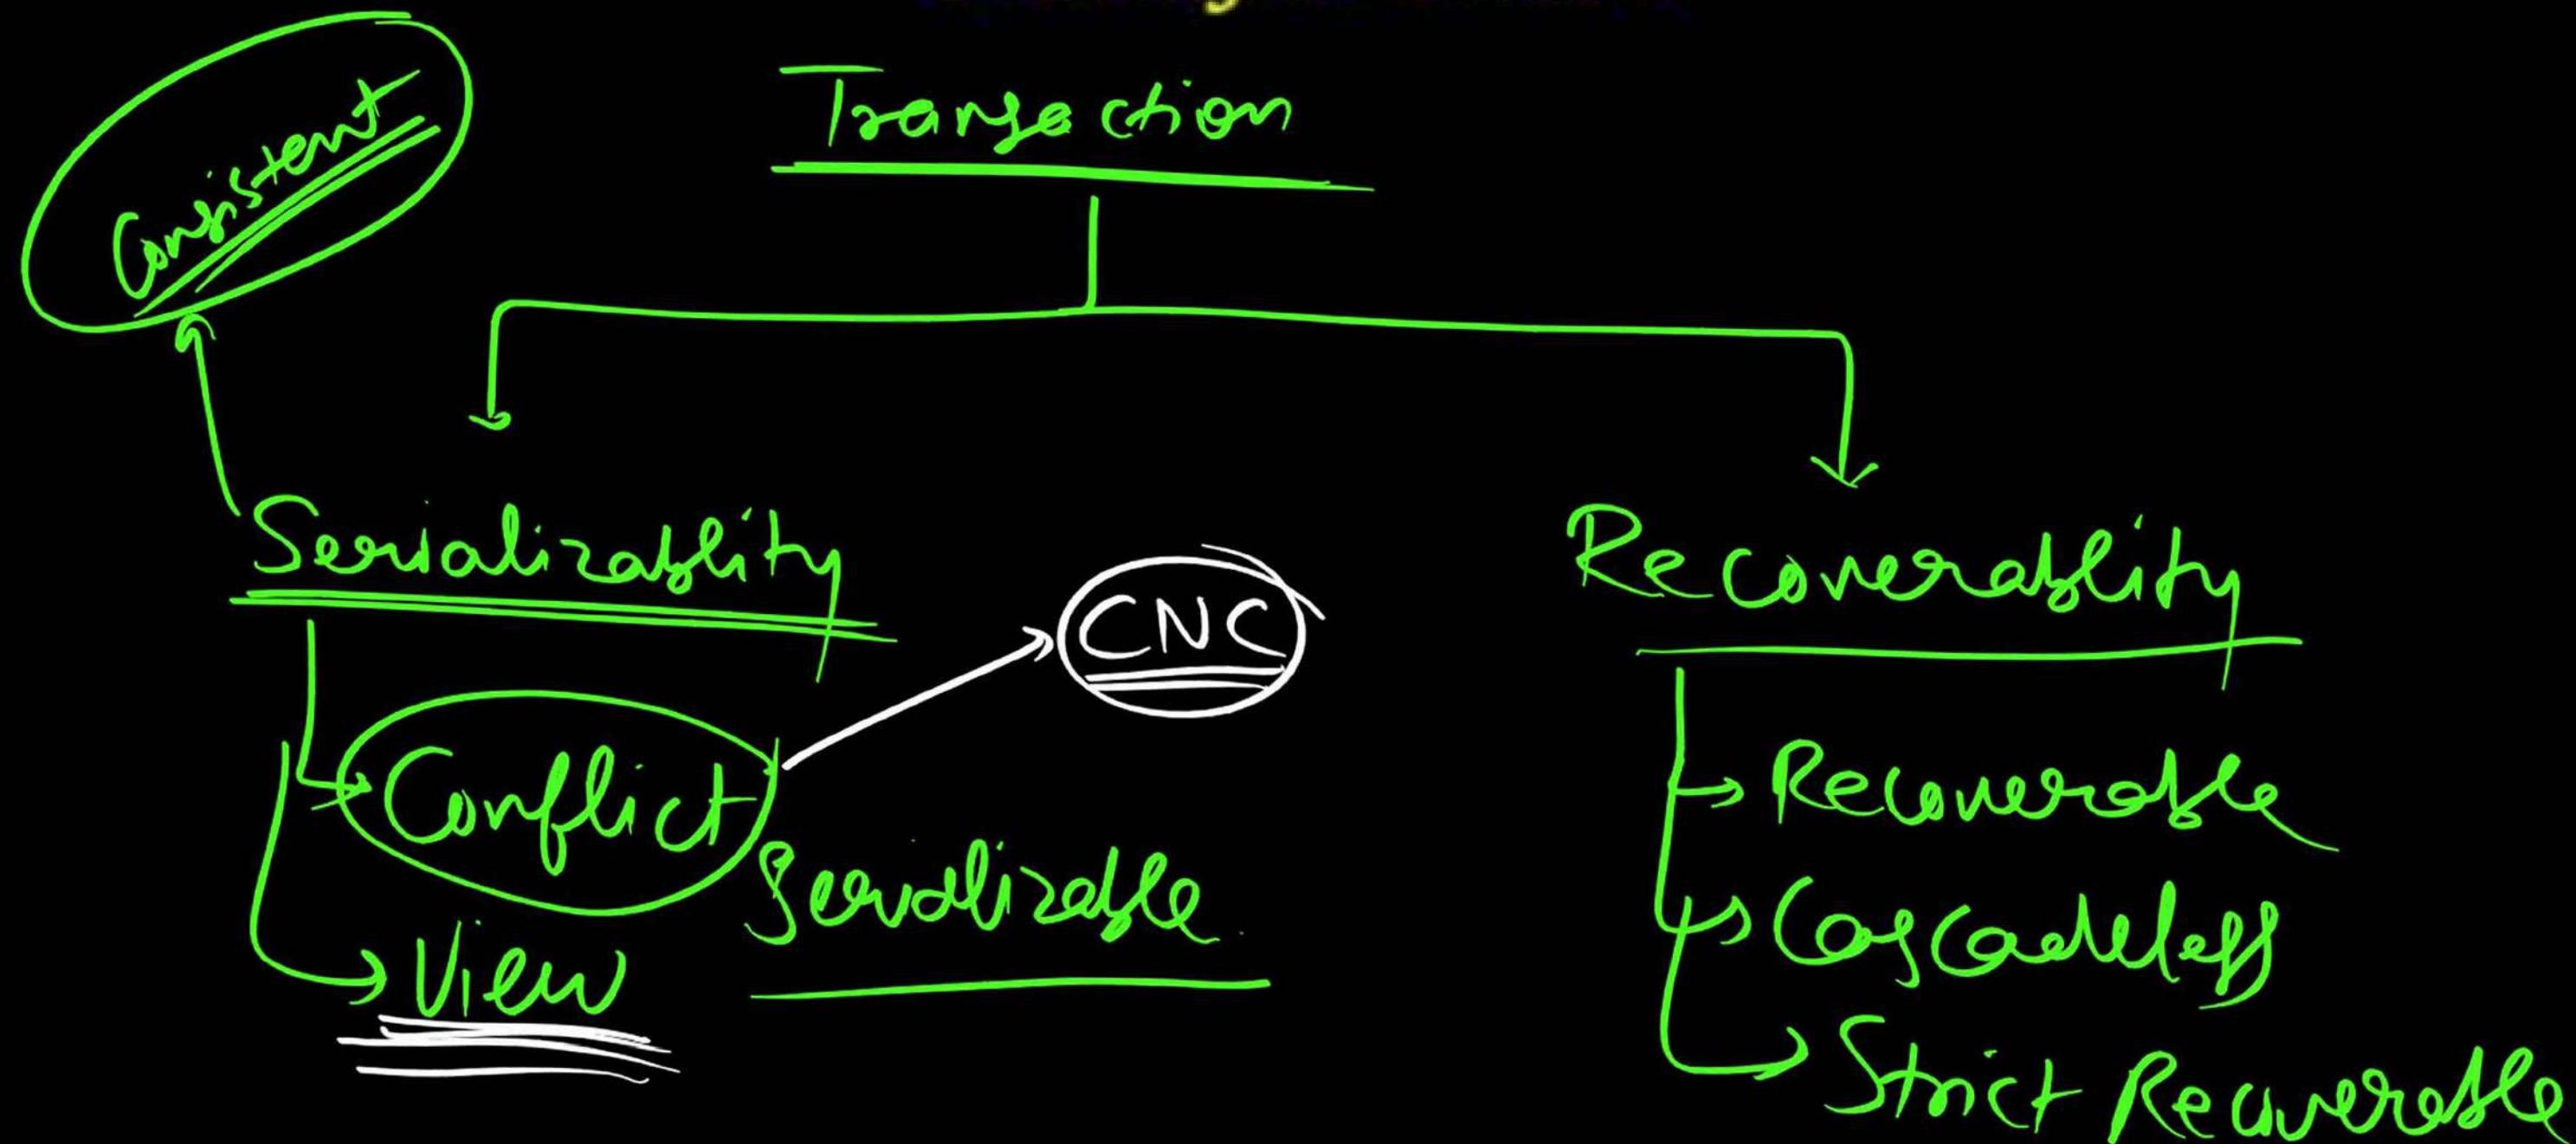
Lecture_05



Vijay Agarwal sir

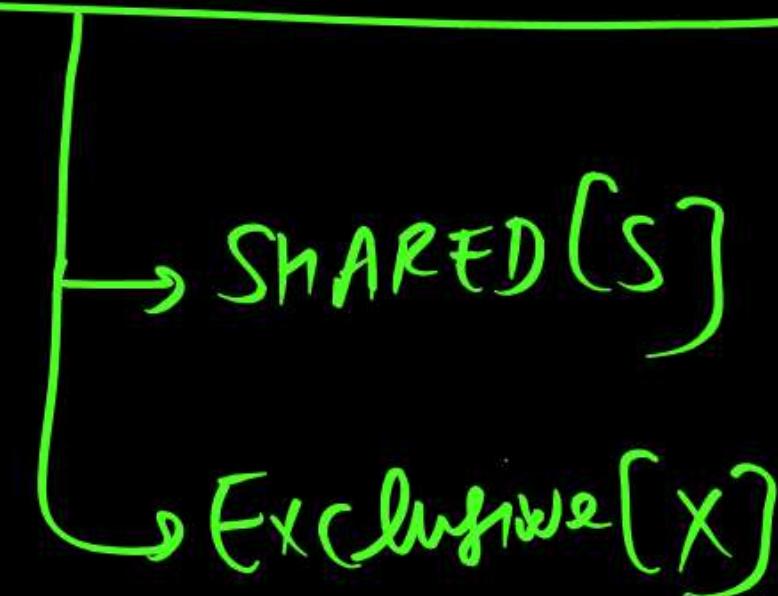


Today's GOAL



Implementation of Concurrency Control

Lock Based Protocol.



SHARED Lock [S]



Only Read

Syntax

Lock - S(A)

Read(A)

Unlock - S(A)

Exclusive lock [X]

Write

Write
Read

Read
Write

Syntax:

Lock - X(A)

Write(A)

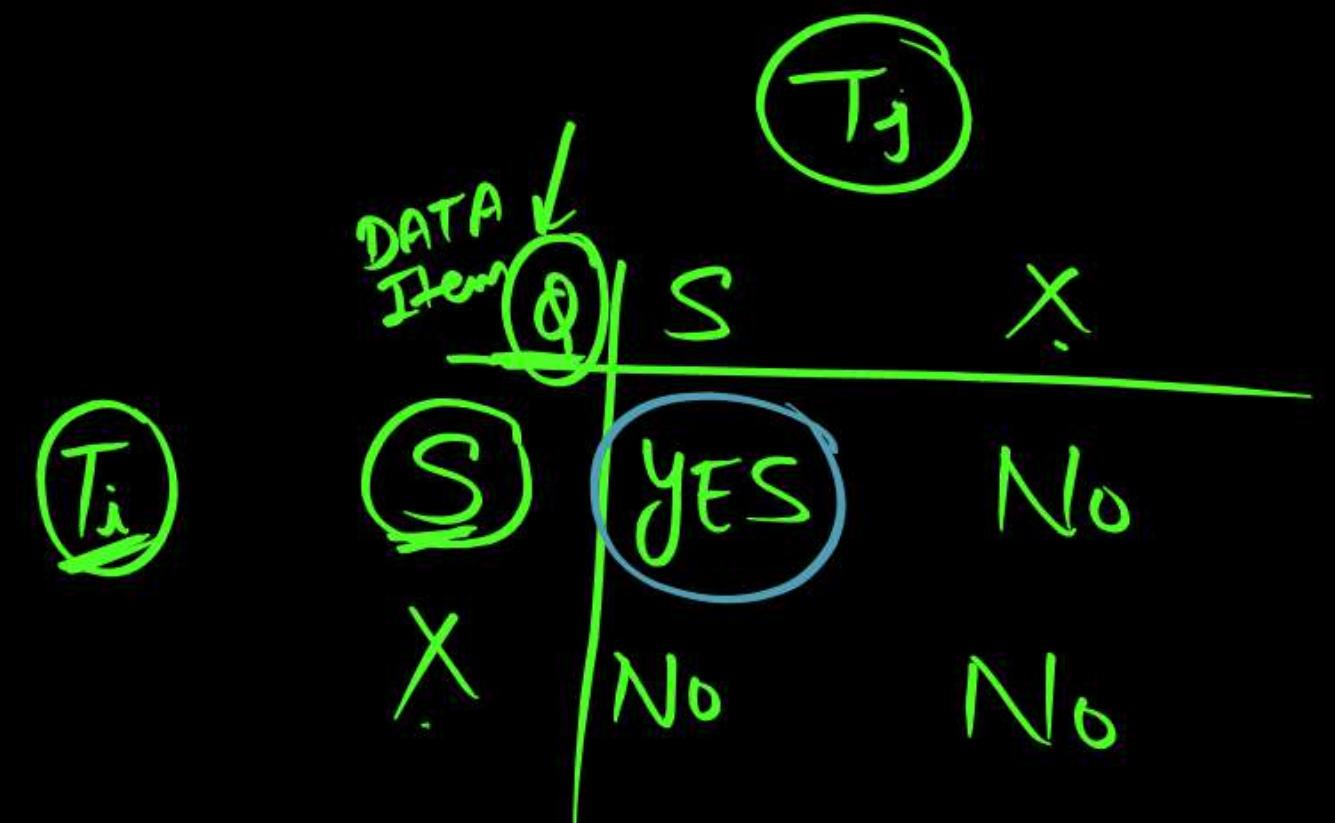
R(A)
WIA

WA
RTA

Unlock - X(A)

Lock-Based Protocols

- ❑ A lock is a mechanism to control concurrent access to a data item
- ❑ Data items can be locked in two modes:
 1. exclusive (X) mode. Data item can be both reads as well as written. X-lock is requested using lock-X instruction.
 2. Shared (S) mode. Data item can only be read. S-lock is requested using lock-S instruction.
- ❑ Lock requests are made to concurrency-control manager. Transaction can proceed only after request is granted.

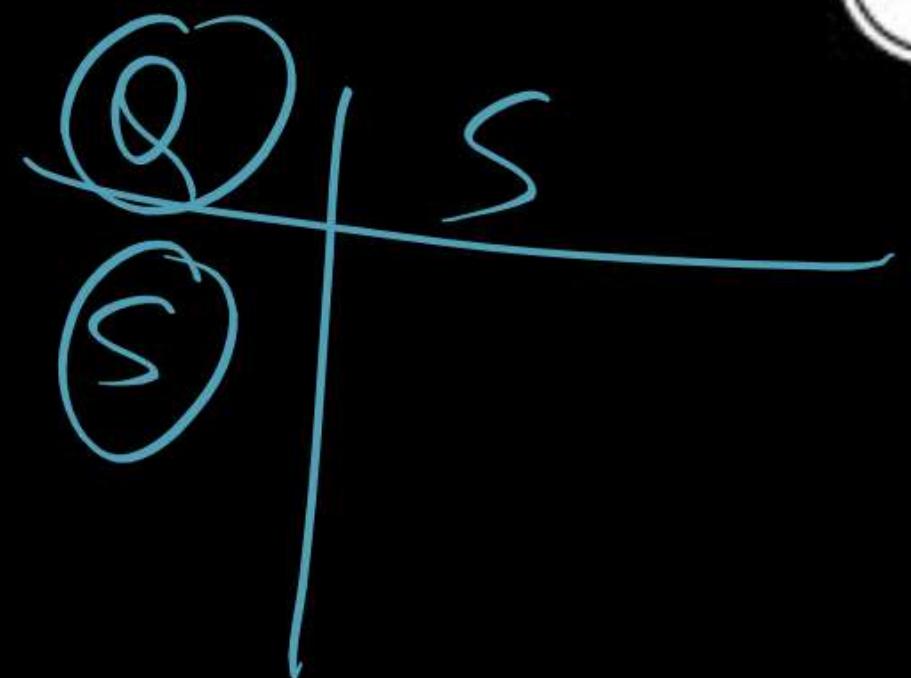


Lock-Based Protocols (Cont.)

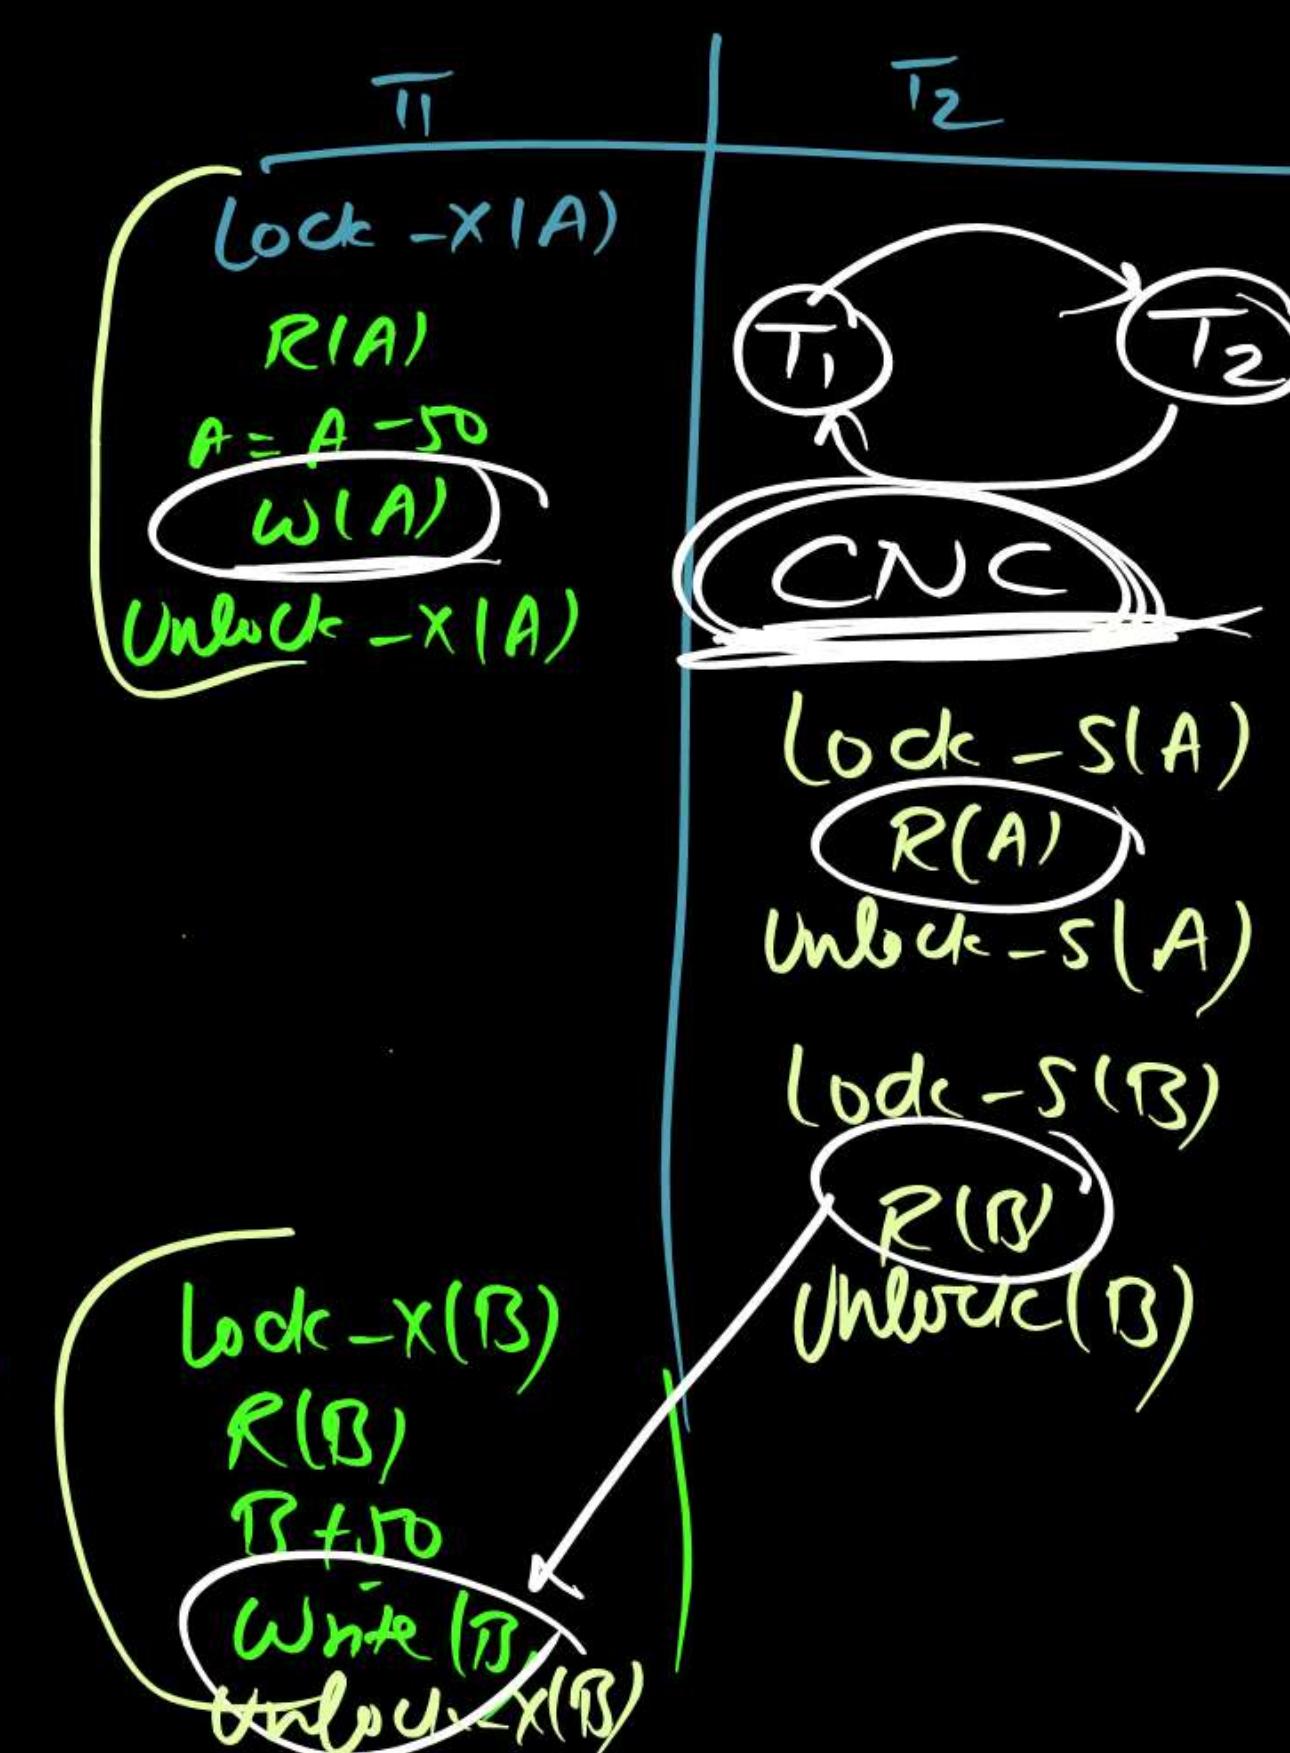
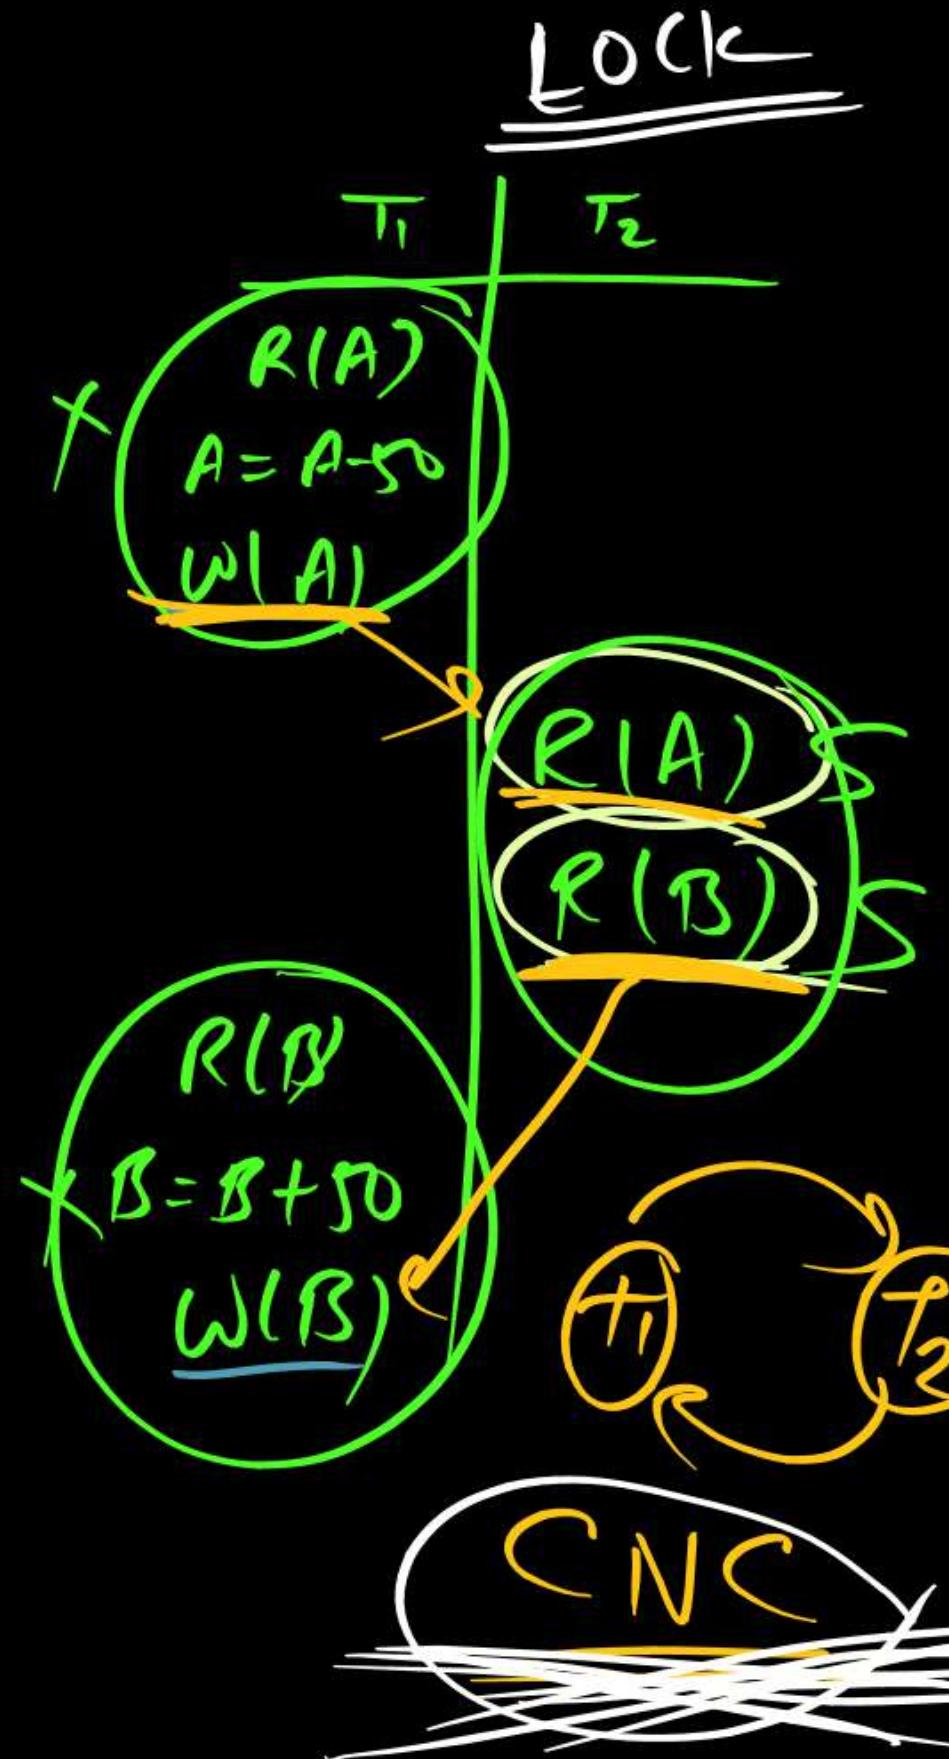
Lock-compatibility matrix

	S	X
S	true	false
X	false	false

- A transaction may be granted a lock on an item if the requested lock is compatible with lock already held on the item by other transactions.
- Any number of transactions can hold shared locks on an item
- But if any transaction holds an exclusive lock on the item no other transaction may hold any lock on the item.



P
W



Grant - $X(A, T_1)$

Release / $-X(A, T_1)$

grant - $S(A, T_2)$

release - $S(A, T_2)$

grant - $S(B, T_2)$

release - $S(B, T_2)$

grant - $X(B, T_1)$

release - $X(B, T_1)$

Schedule with Lock Grants

- A locking protocol is a set of rules followed by all transactions while requesting and releasing locks.
- Locking protocols enforce serializability by restricting the set of possible schedules.

T ₁	T ₂	Concurrency-control manager
lock-X(B) read(B) B:=B - 50 write(B) unlock(B)	lock-S(A, T ₂) read(A) unlock(A) lock-S(B) read(B) unlock(B) display(A+B)	grant-X(B, T ₁) grant-S(A, T ₂) grant-S(B, T ₂) grant-X(A, T ₁)
	σ	

2PL [2 Phase Locking Protocol]

Each Transaction Issue lock & unlock Request in
Two Phases

- ① Growing Phase [Acquire the lock but must NOT release Any lock]
- ② Shrinking Phase [Release the lock but MUST NOT ask for Any New lock]

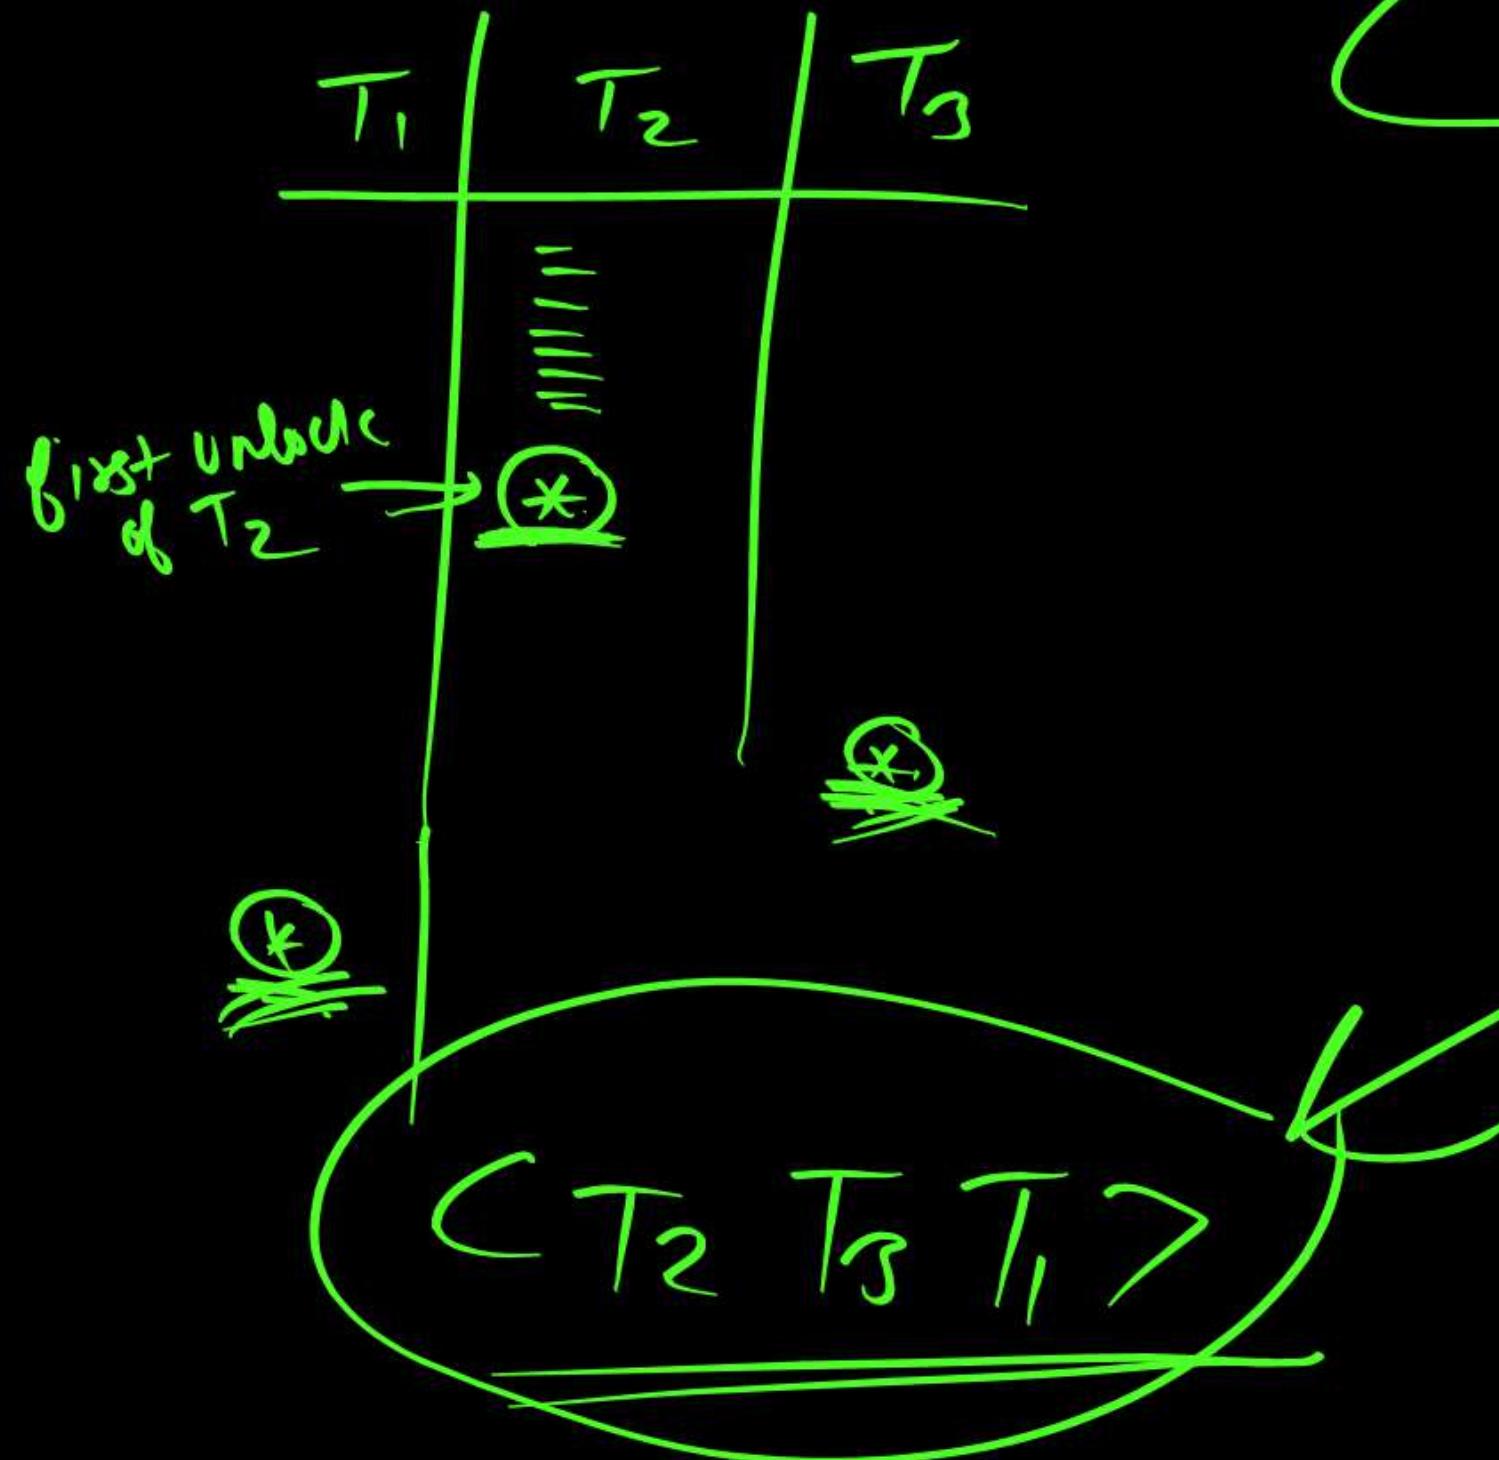
Note

Each Transaction finish first its Growing

Phase then Shrinking Phase

LOCK Point : from where Shrinking Phase Start .

Position of last lock @ First Unlock operation.

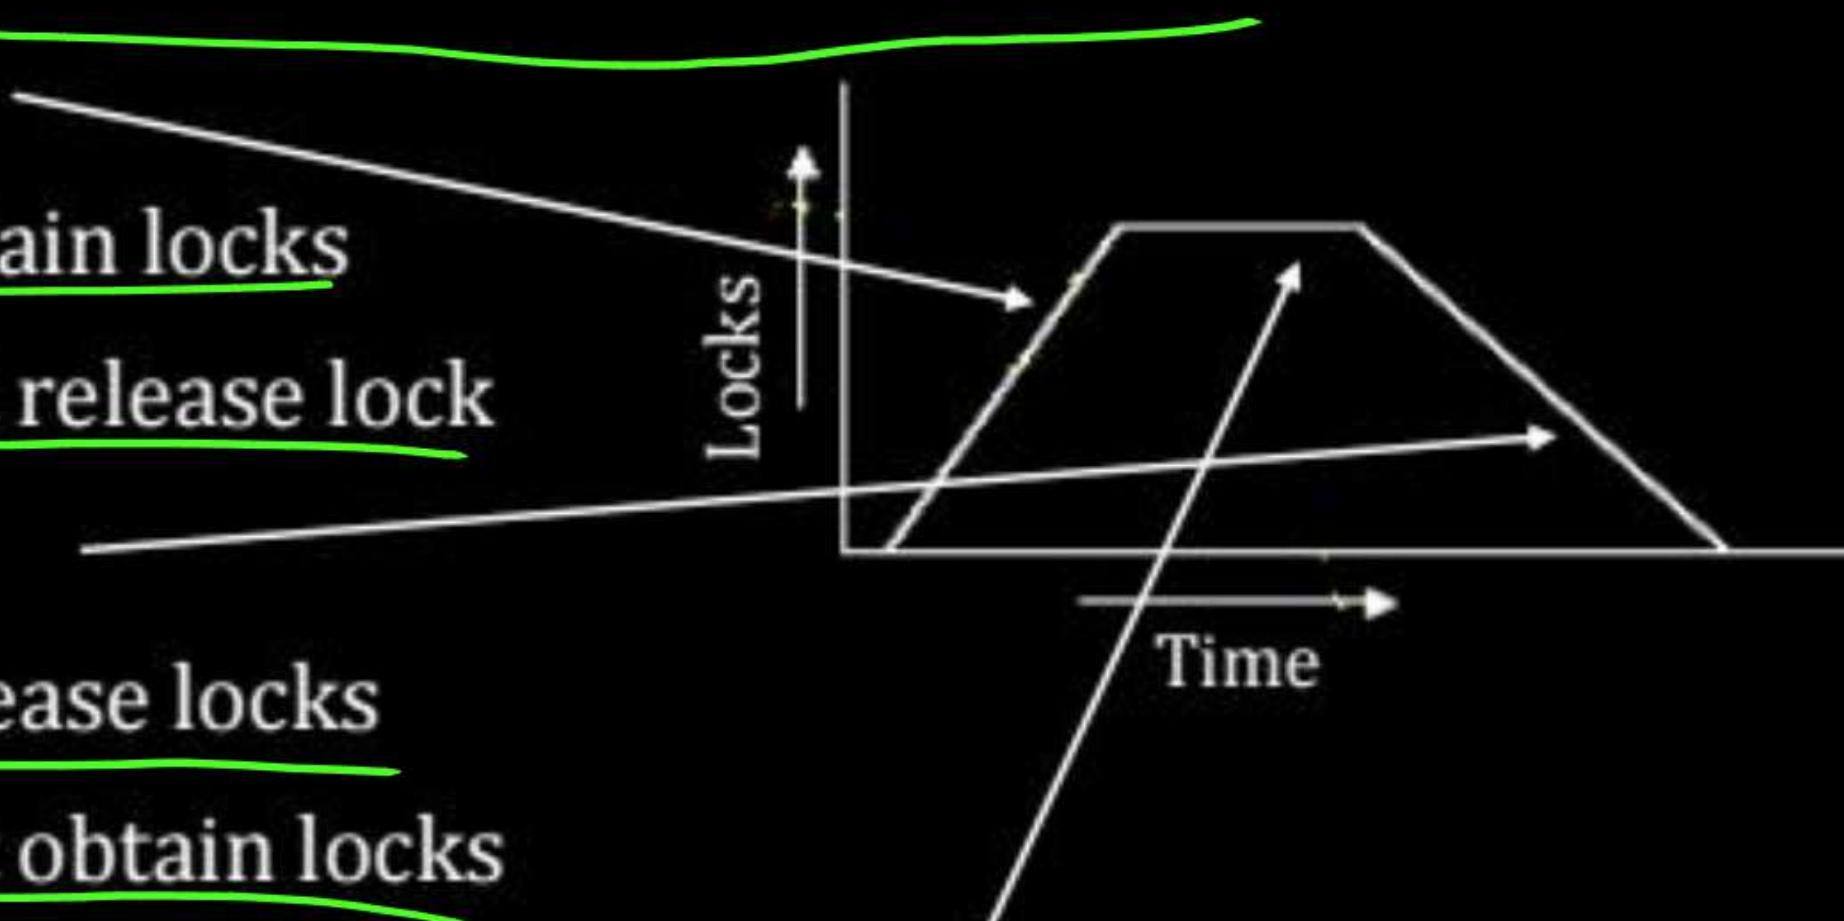


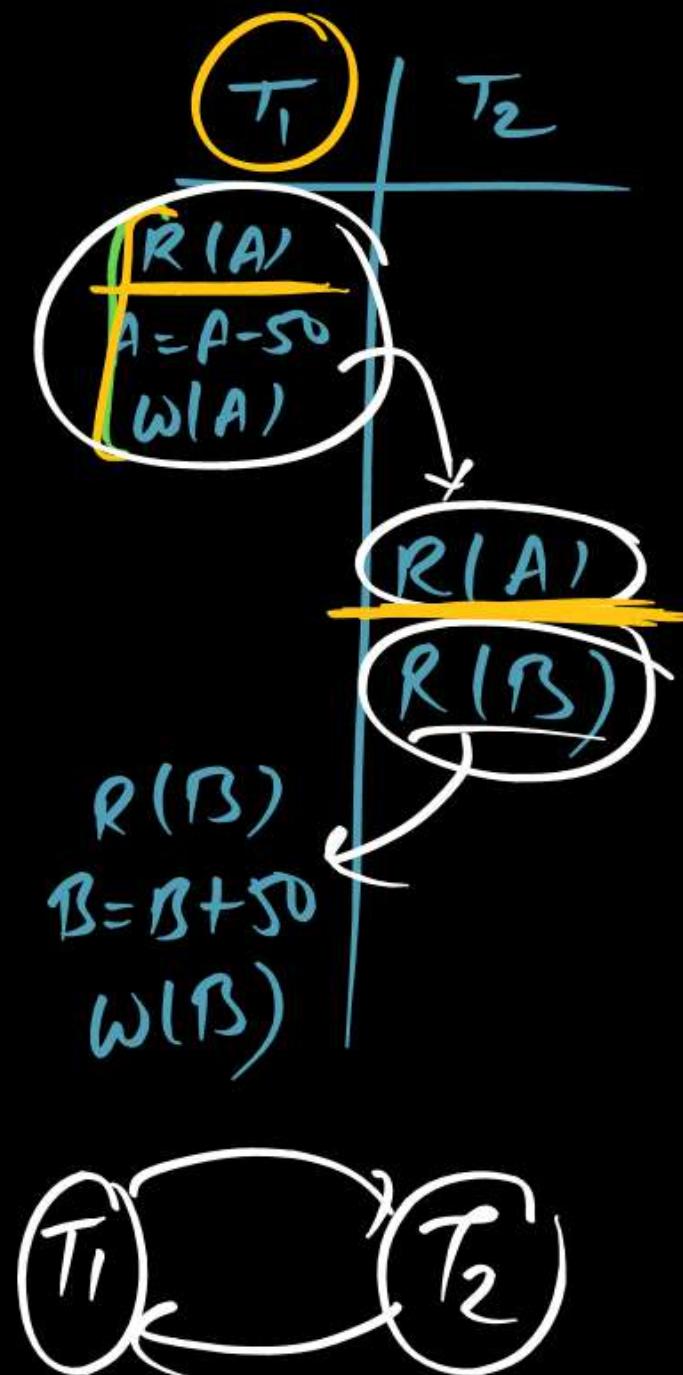
LOCK Point

↳ Serializability Order

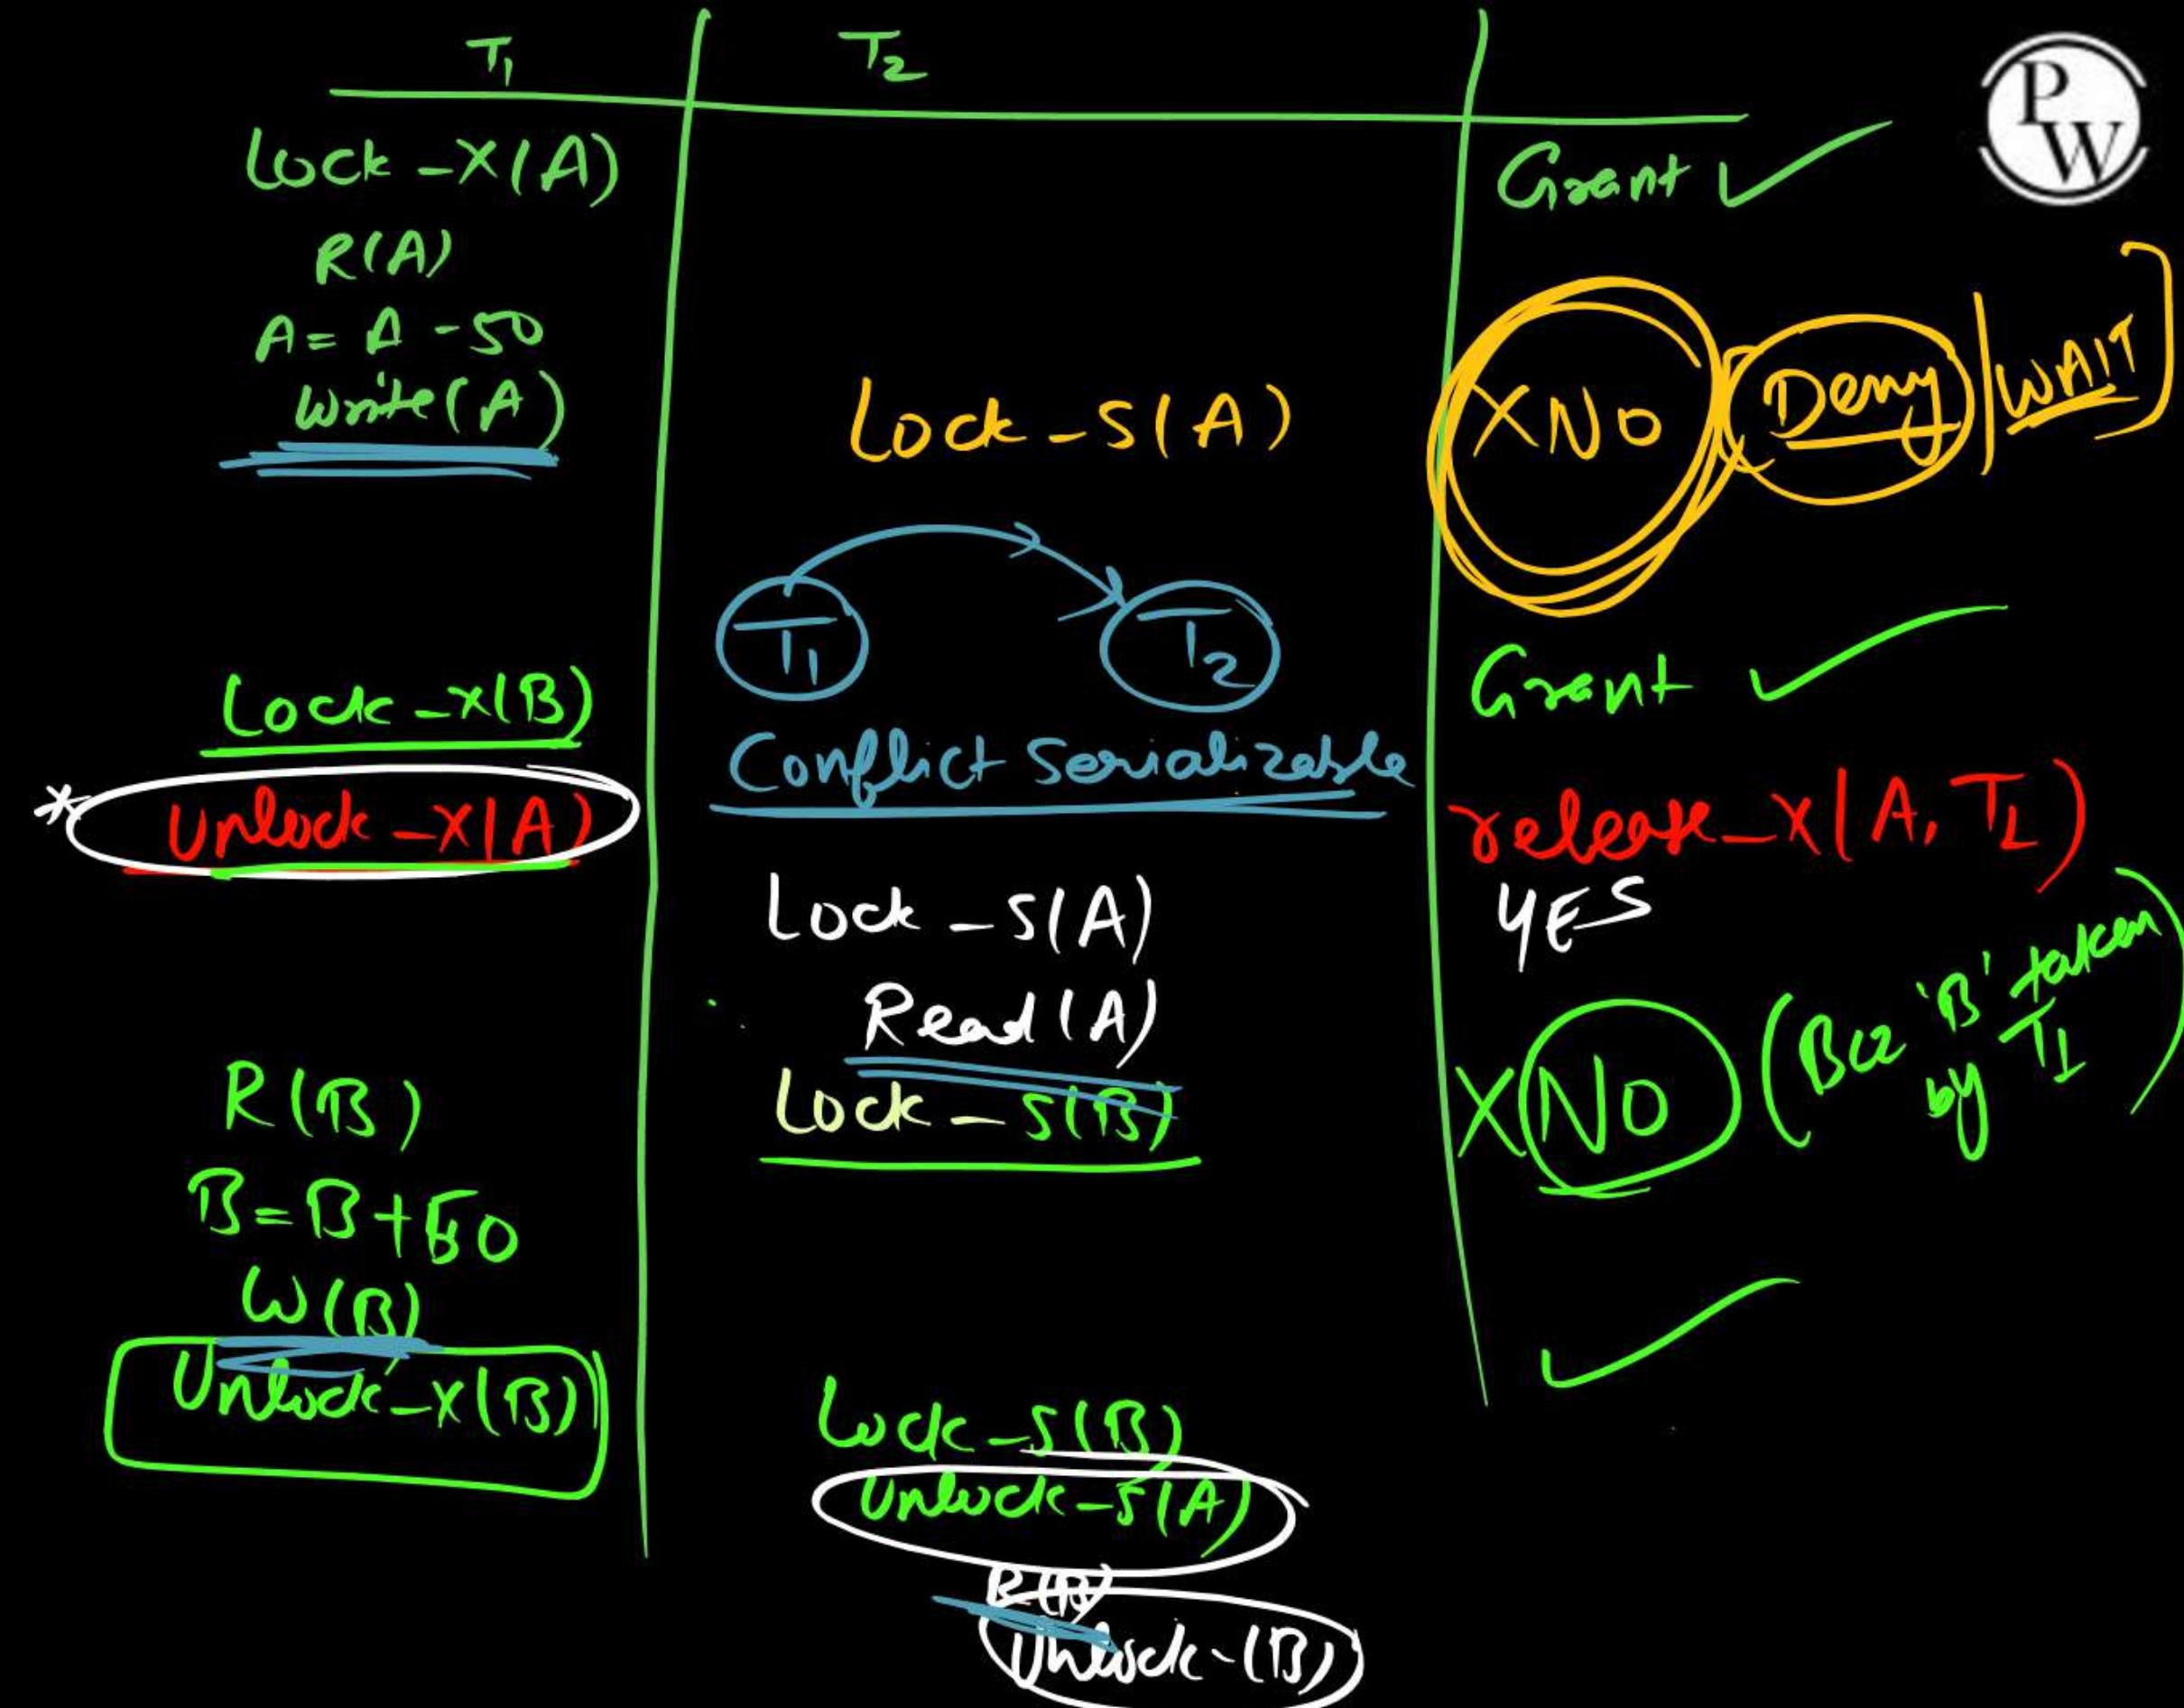
The Two-Phase Locking Protocol

- A protocol which ensures conflict-serializable schedules.
- Phase 1: Growing Phase
 - ❖ Transaction may obtain locks
 - ❖ Transaction may not release lock
- Phase 2: Shrinking Phase
 - ❖ Transaction may release locks
 - ❖ Transaction may not obtain locks
- The protocol assures serializability. It can be proved that transactions can be serialized in the order of their lock points (i.e., the point where a transaction acquired its final lock).





Cycle Not Conflict.



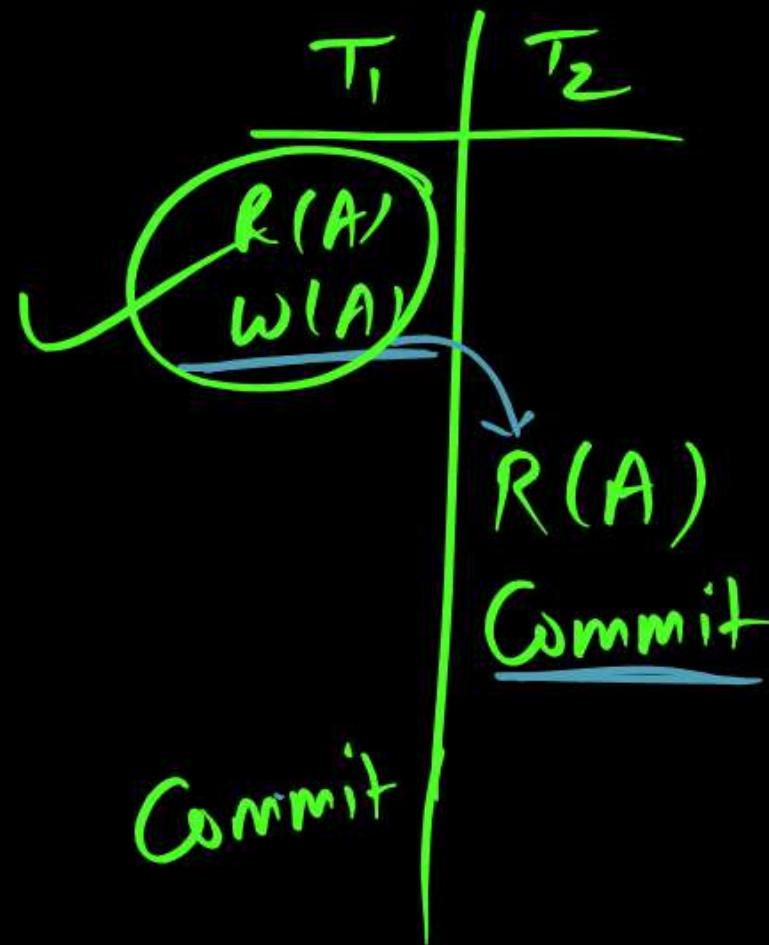
Note

If a Schedule followed by 2PL (Applied 2PL successfully)

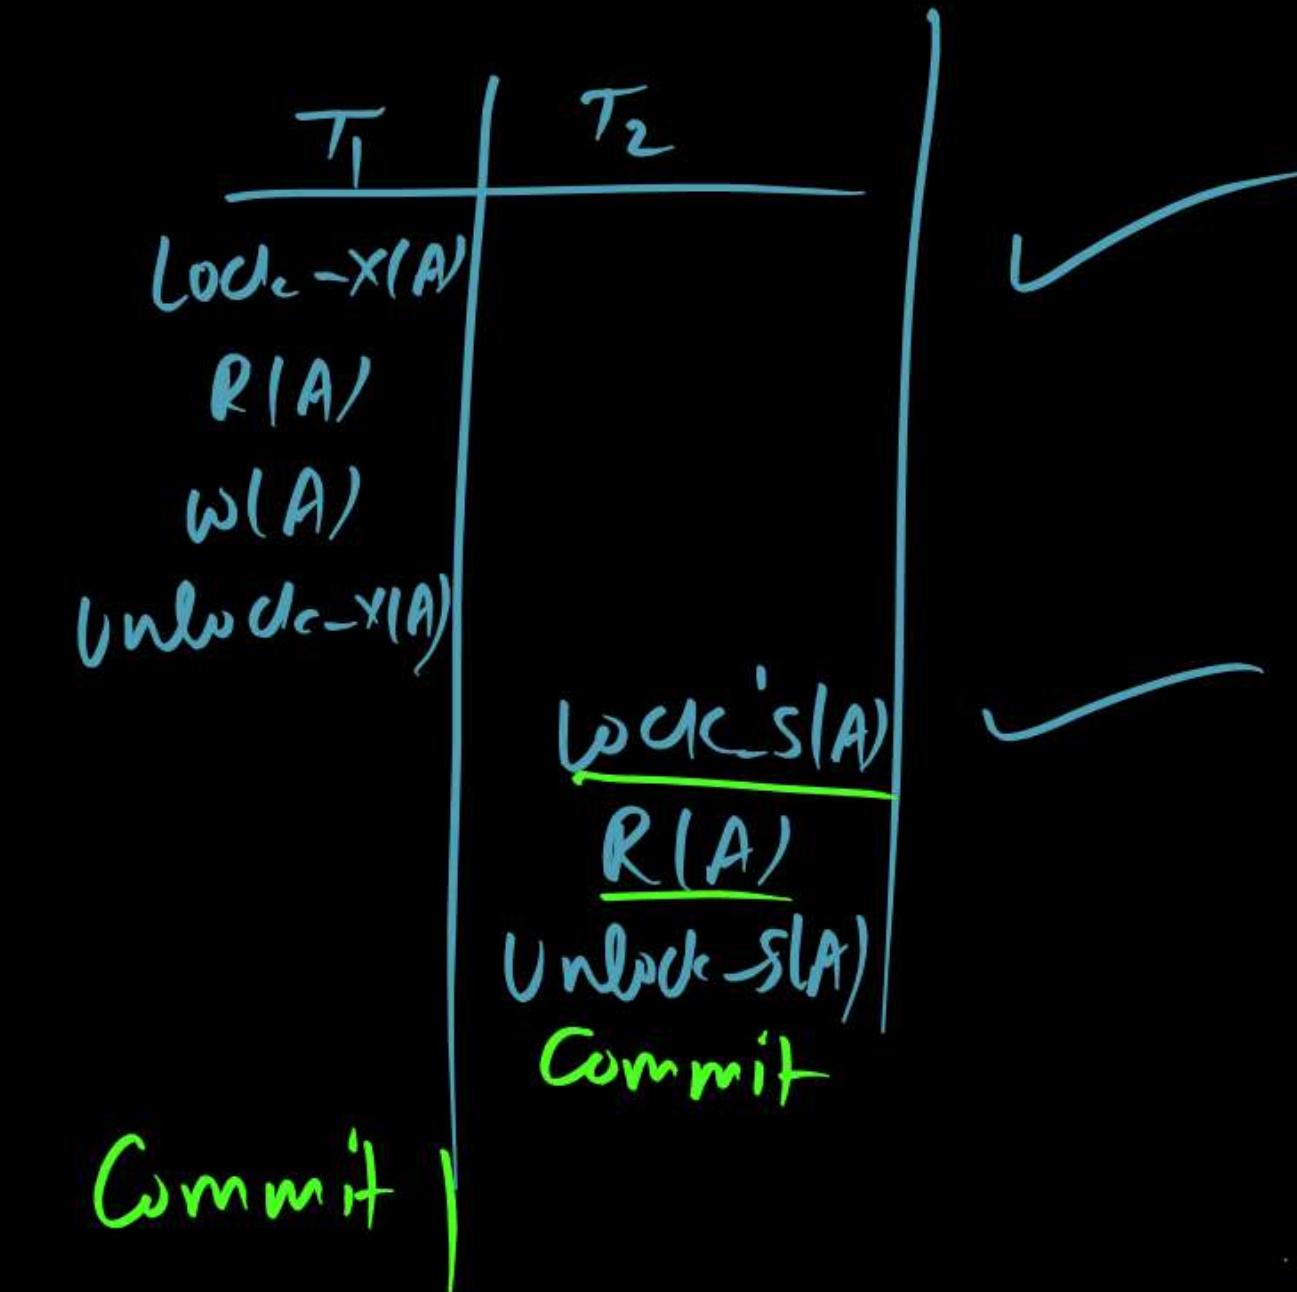
then its Conflict Serializable Schedule

Note

2PL ensure Serializability

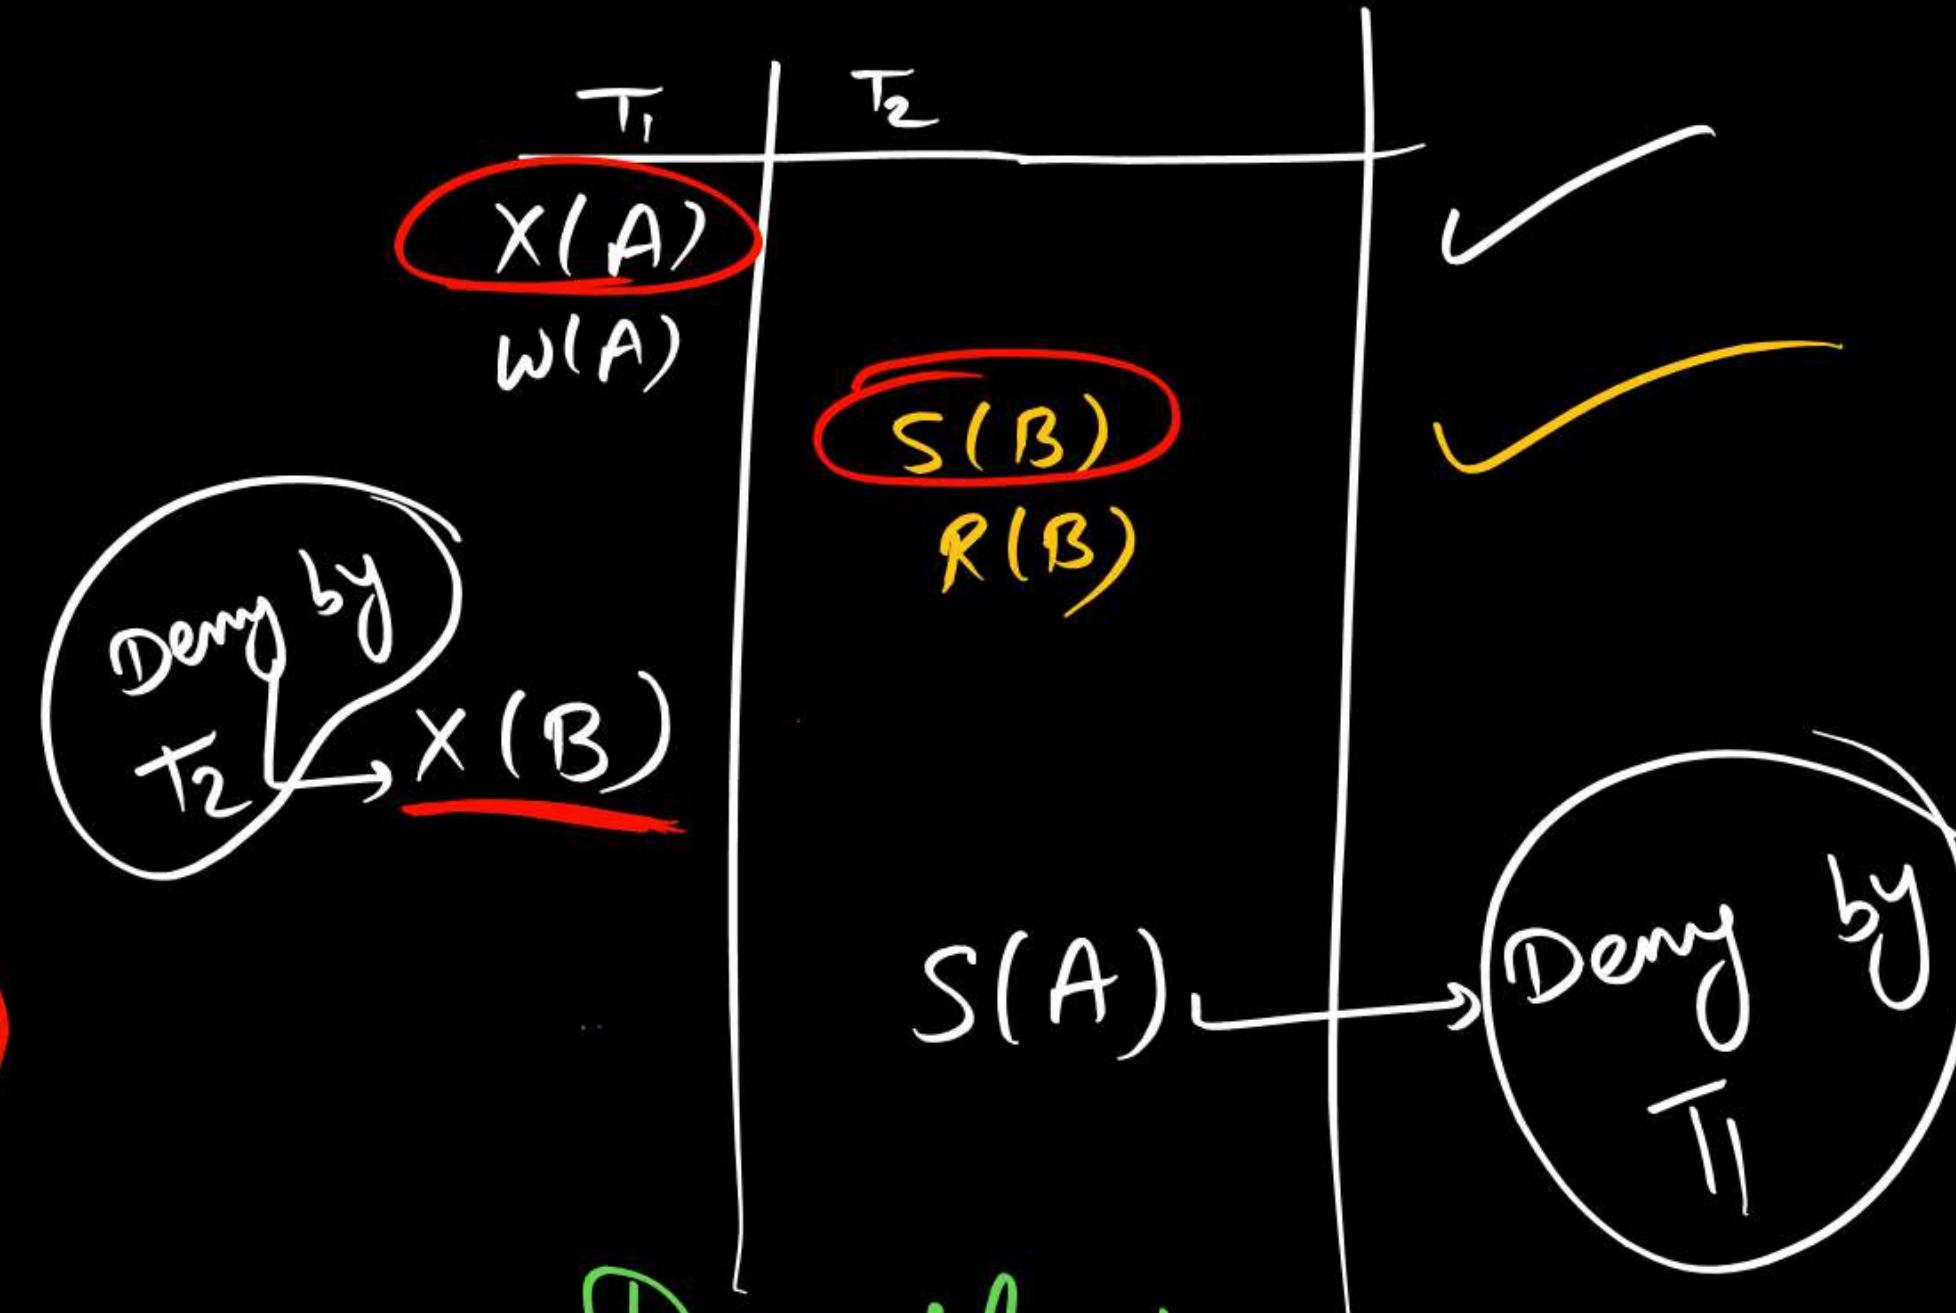
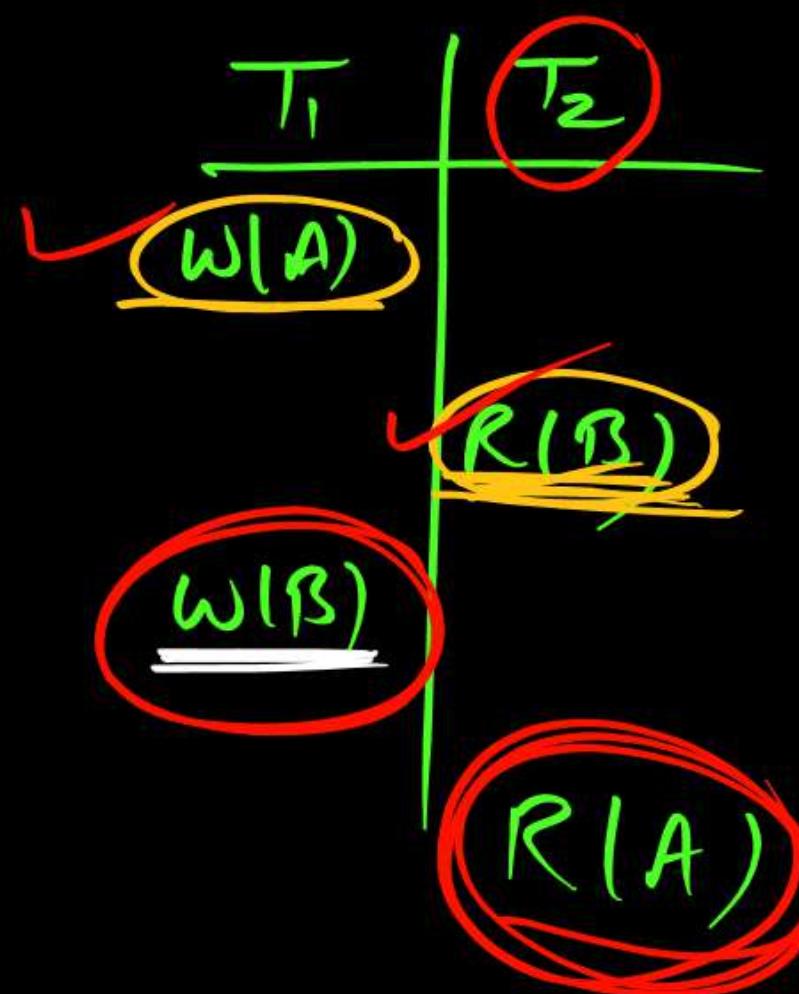


Irrecoverable

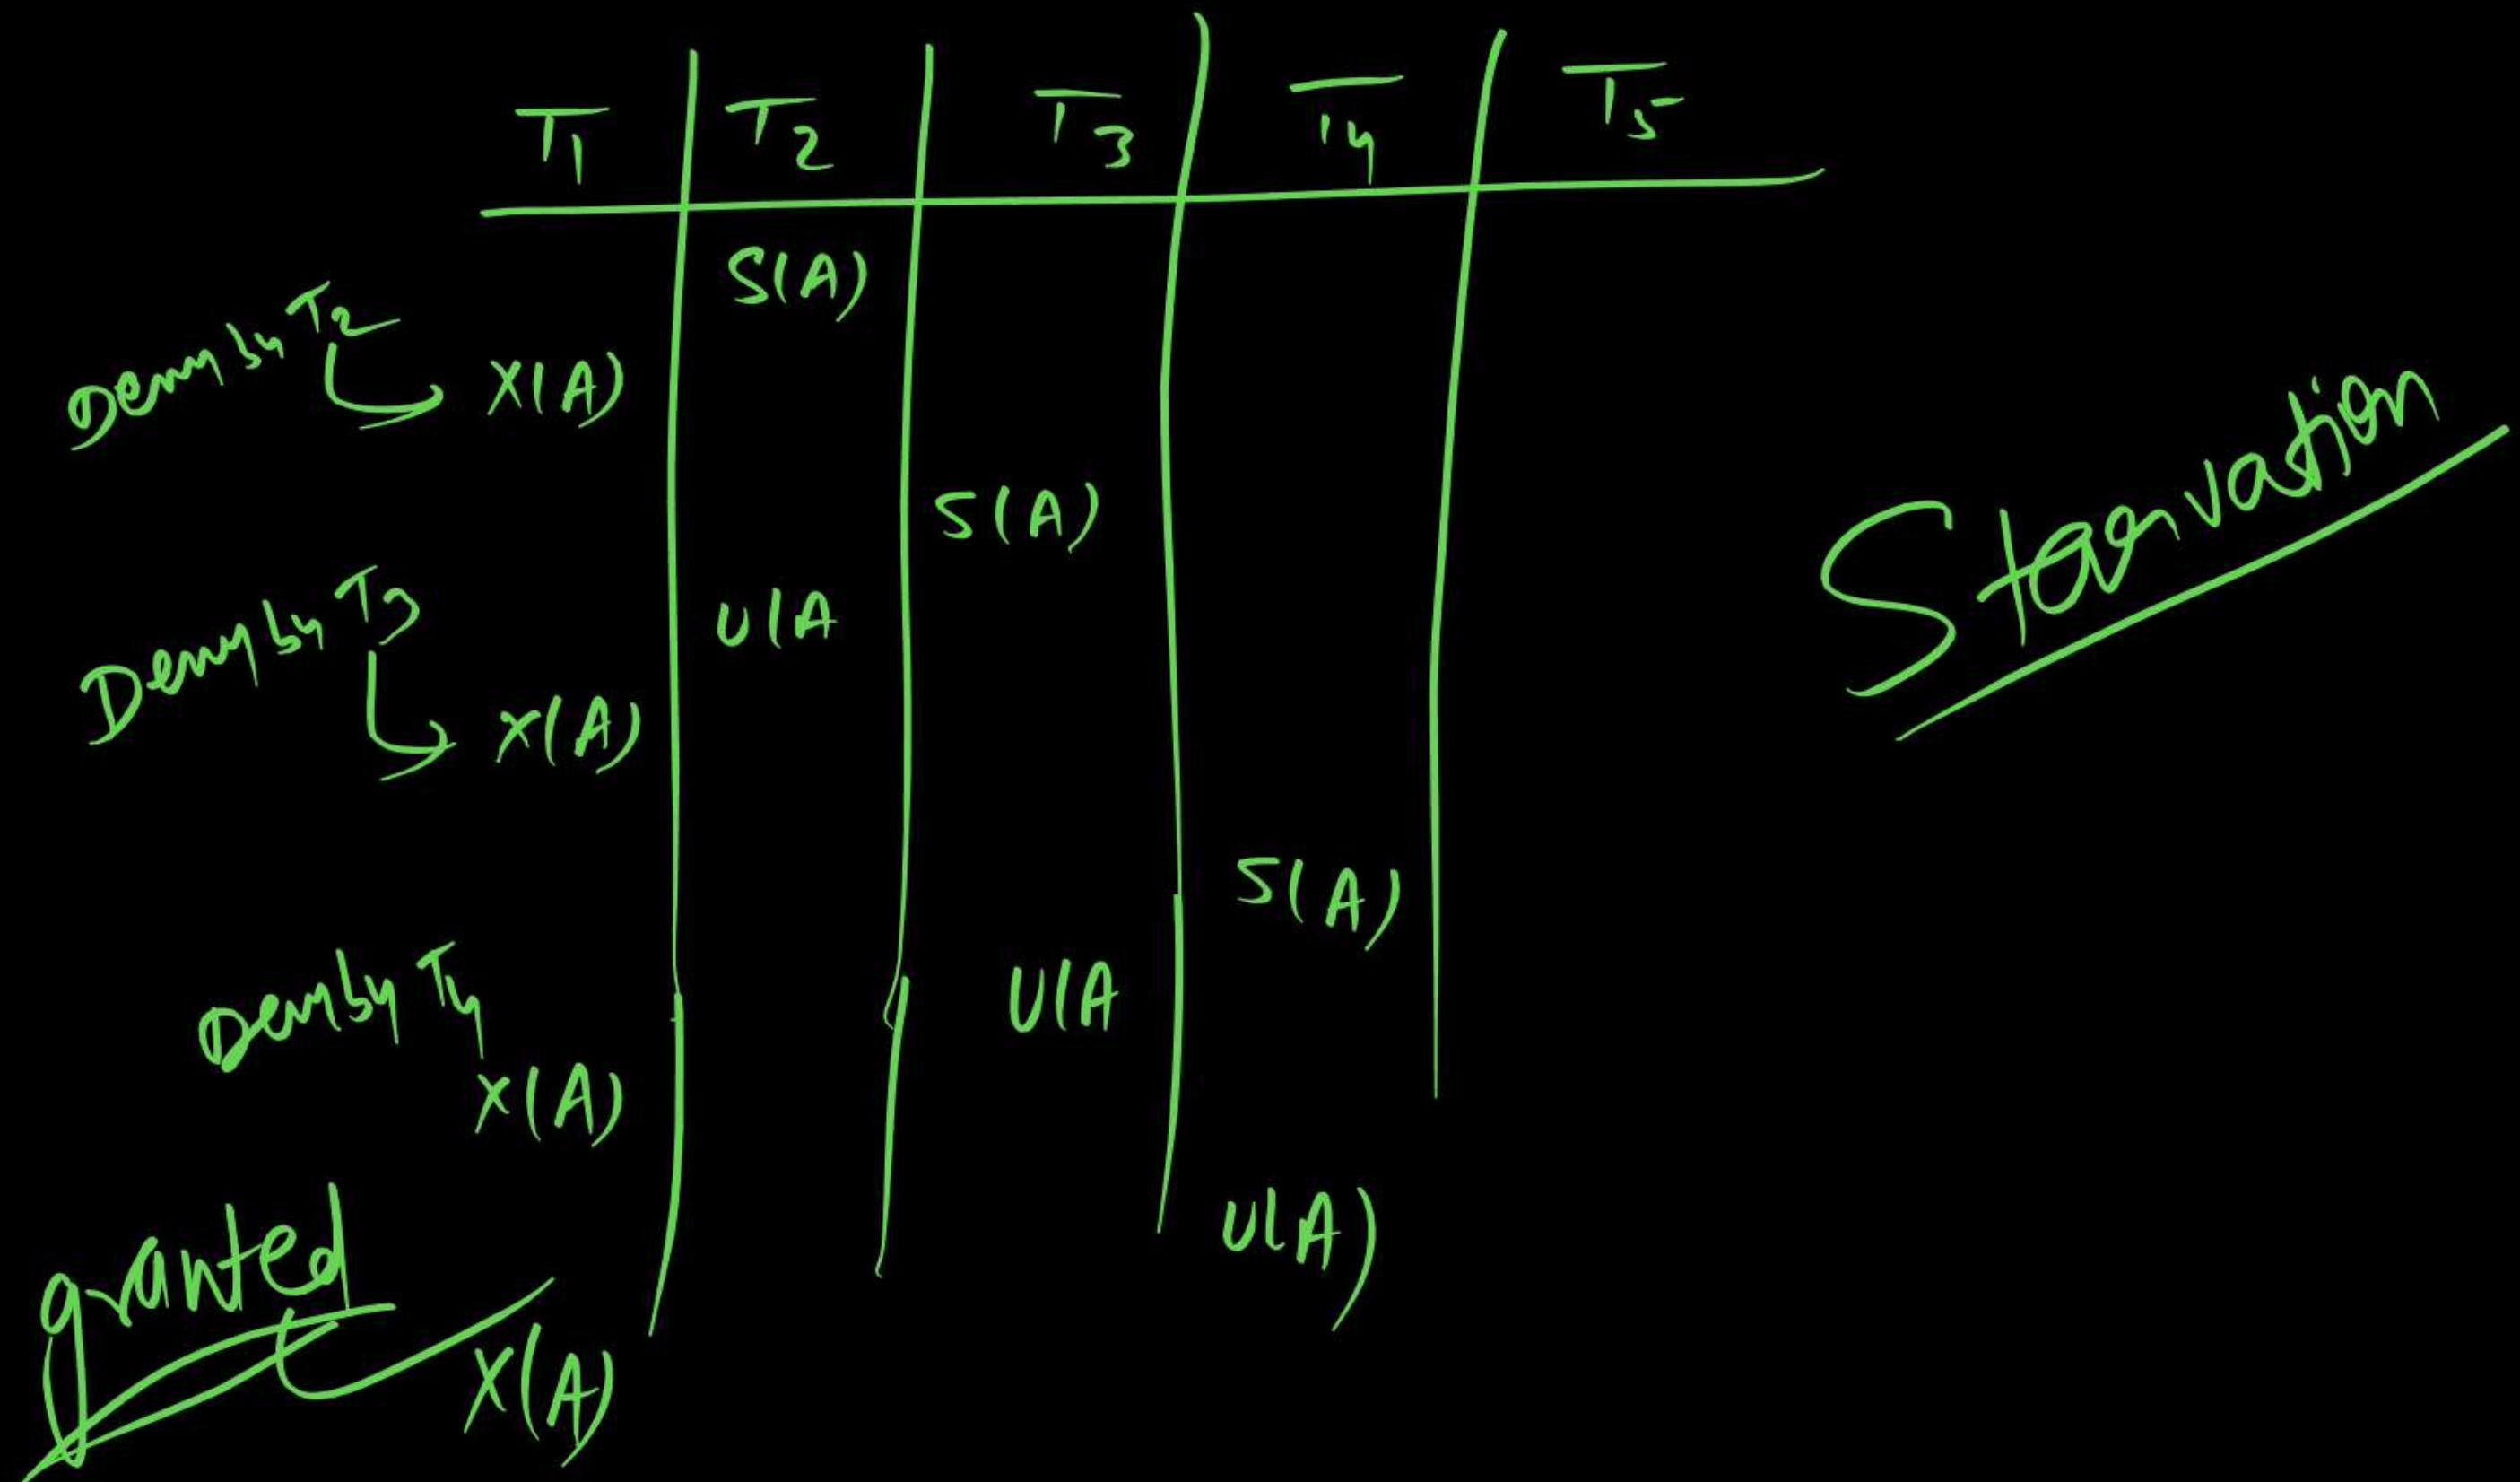


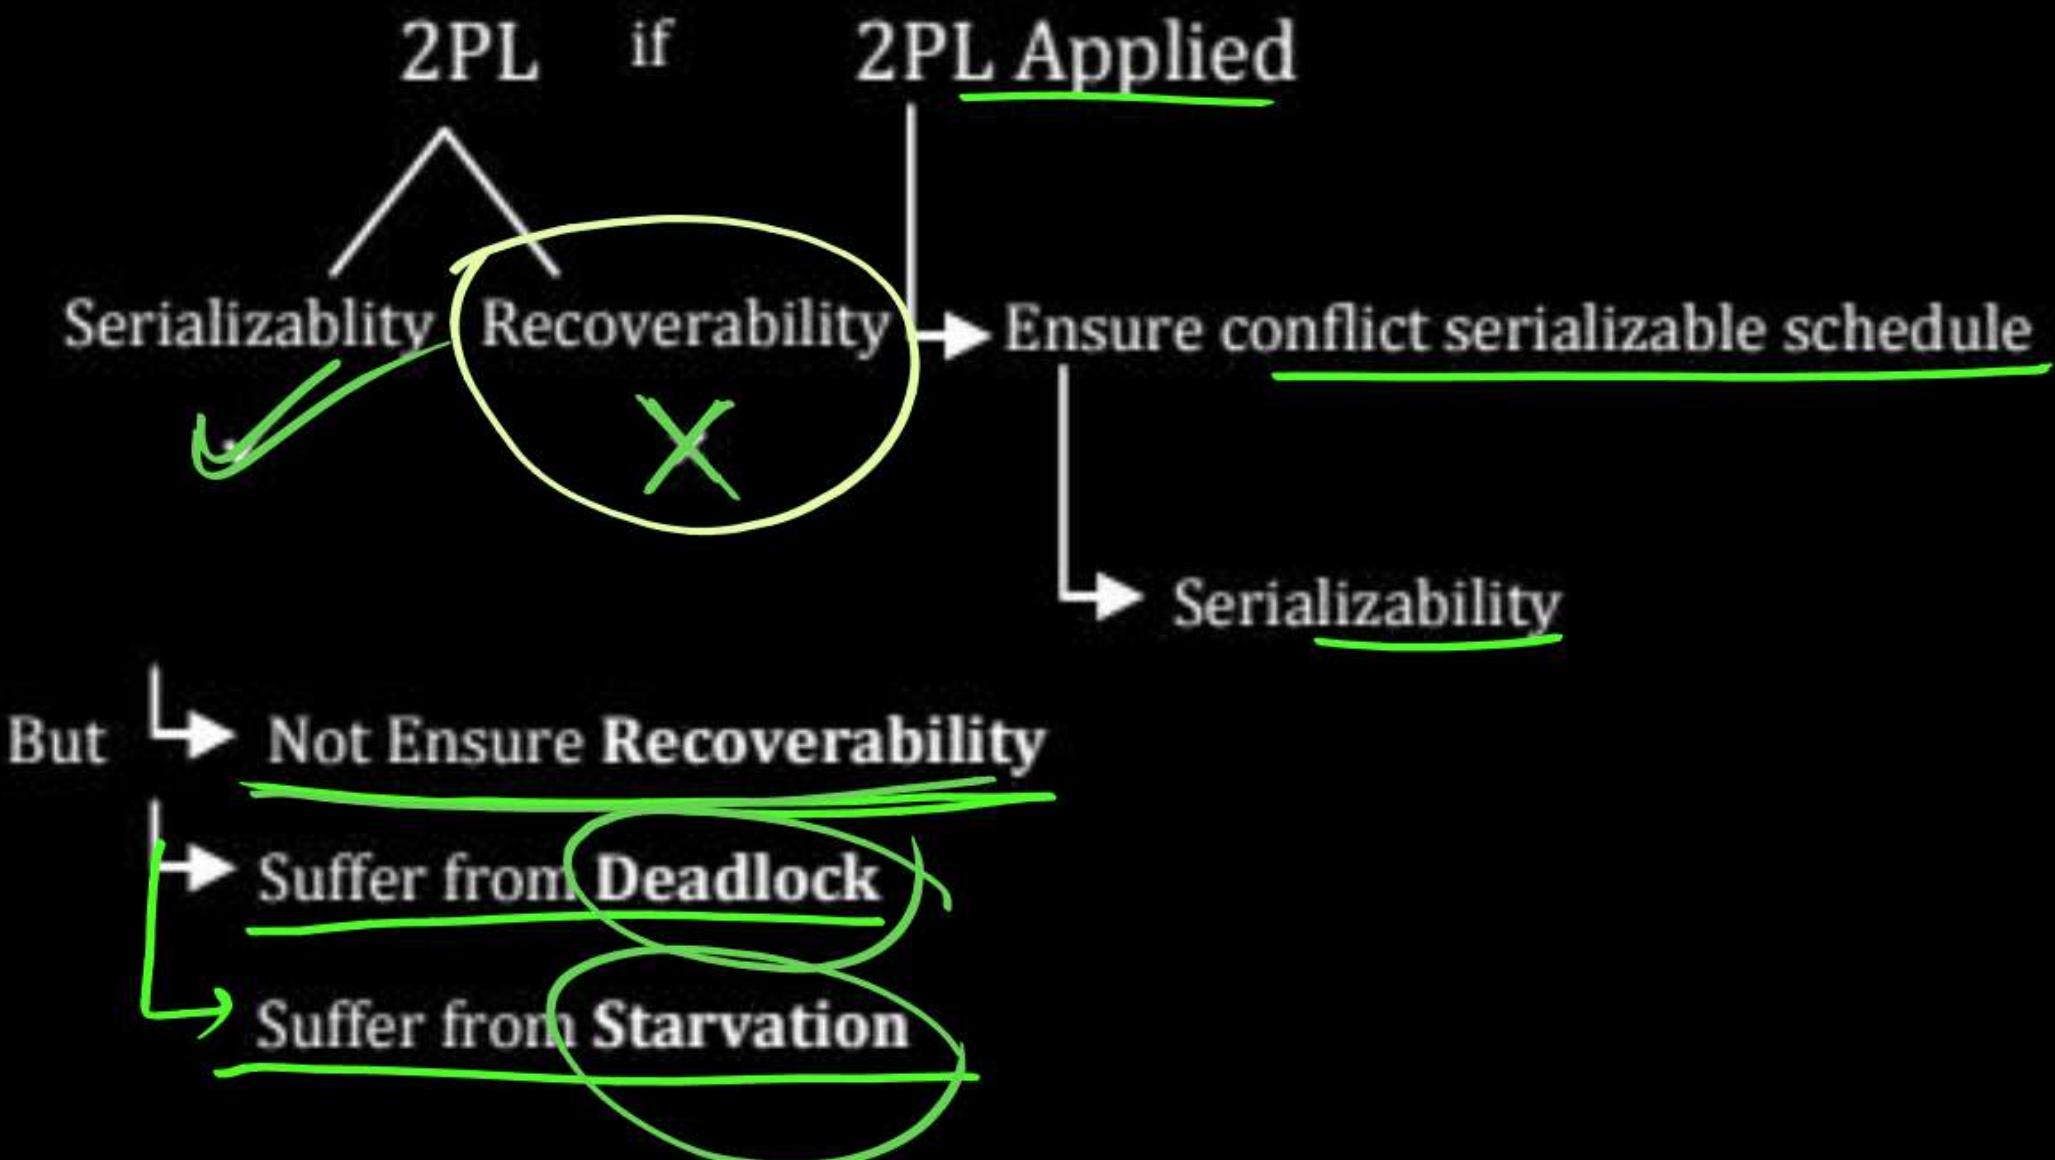
Its Irrecoverable but followed by QPC.

P
W



Deadlock

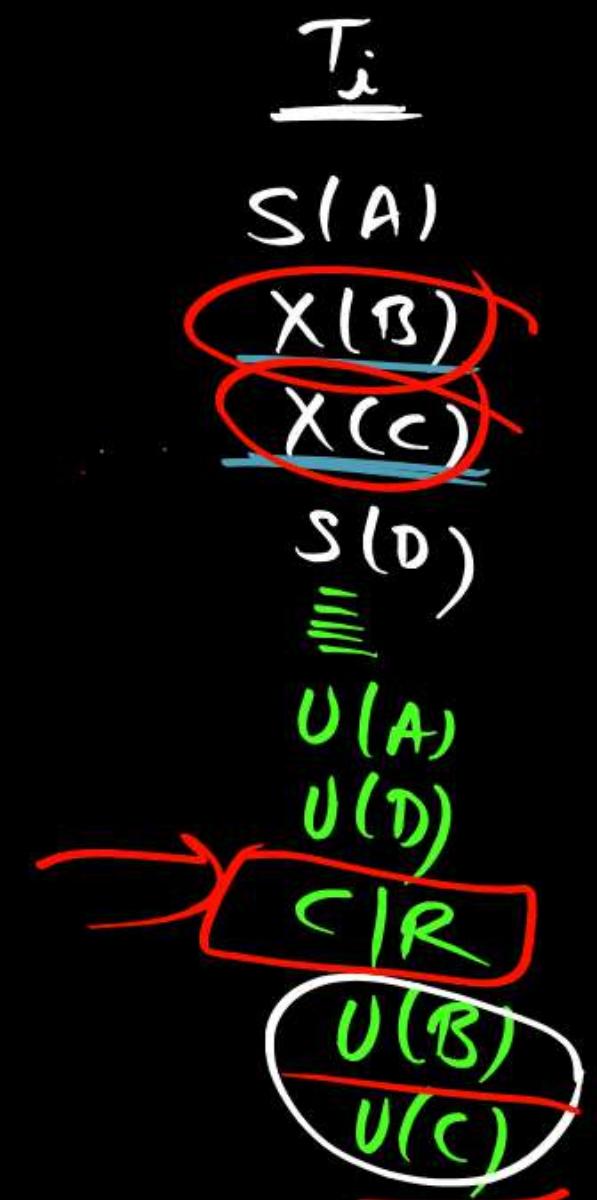


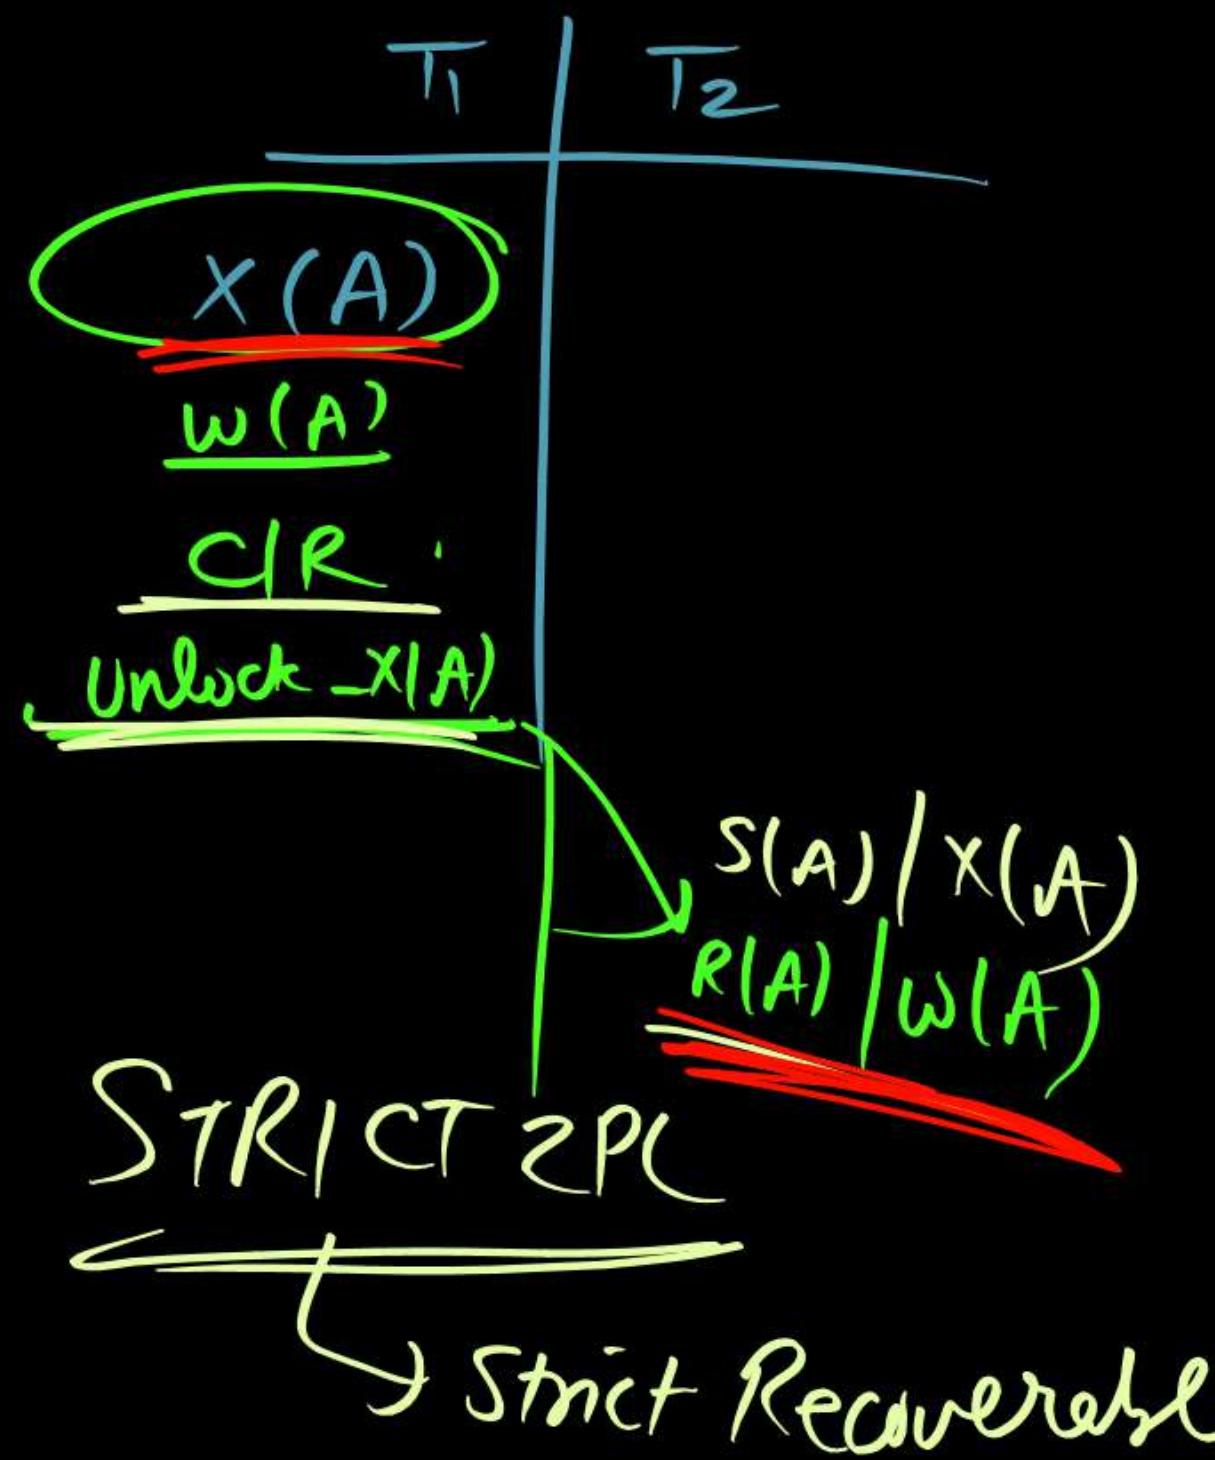


STRICT 2PL

2PL + All Exclusive lock must be held
Until Commit | Rollback of the Transaction

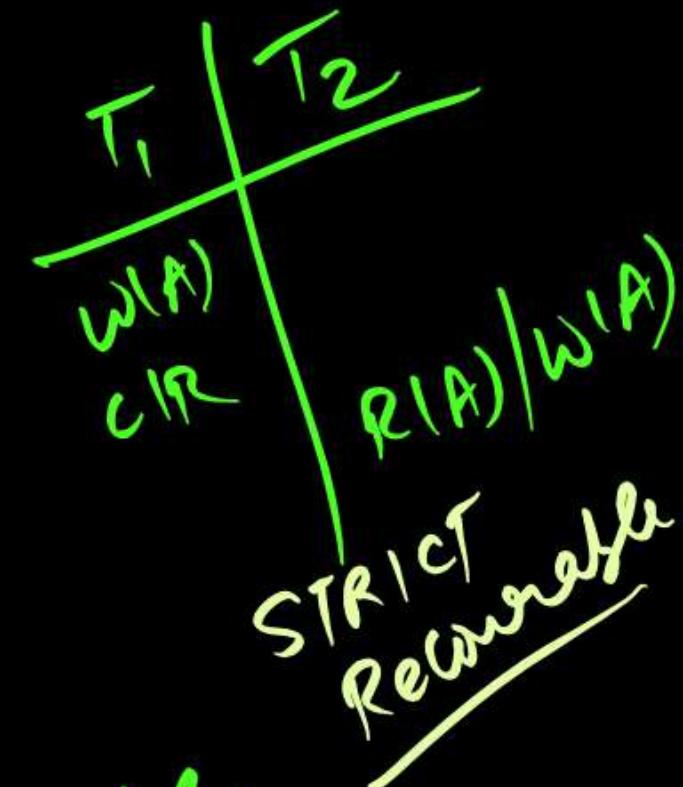
OR
2PL + All X lock Release
After Commit | Rollback of the Transaction





STRICT 2PL

- Recoverable
- Cascadless
- Strict Recoverable



(S/X)

Rigorous 2PC

• 2PC + All Lock Hold Until Commit/Rollback

⑥ ↓

S(A)
X(B)
X(C)
S(D)2PC + Release All lock After C/R

• Deadlock

Starvation

C/RU(A)
U(B)
U(C)
U(D)

STRICK 2PL
RIGOROUS 2PL

→ Subtle Deadlock
&
Starvation

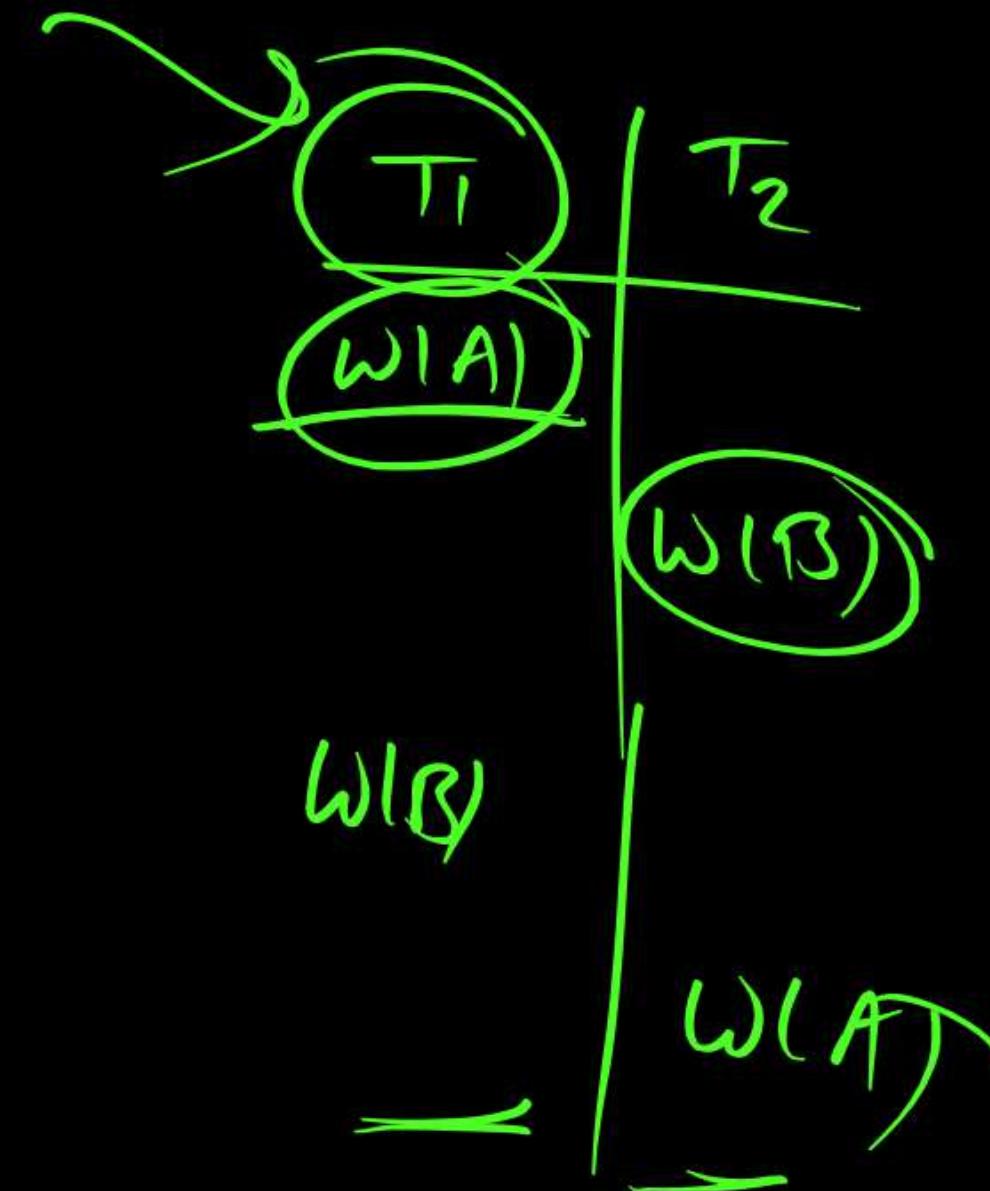
The Two-Phase Locking Protocol (Cont.)

- ❑ Two-phase locking does not ensure freedom from deadlocks
- ❑ Extensions to basic two-phase locking needed to ensure recoverability of freedom from cascading roll-back
 - ❖ Strict two-phase locking: a transaction must hold all its exclusive locks till it commits/aborts.
 - Ensures recoverability and avoids cascading roll-backs
 - ❖ Rigorous two-phase locking: a transaction must hold all locks till commit/abort.
 - Transactions can be serialized in the order in which they commit.
- ❑ Most databases implement rigorous two-phase locking, but refer to it as simply two-phase locking

Conservative 2PL : Acquire all the lock whenever

transaction start &

Release all lock After Commit.



↳ No Deadlock

Only Starvation

1	2	3	4	5
Lock-S(A)	Lock-S(A)	Lock-S(A)	Lock-S(A)	Lock-S(A)
R(A)	R(A)	R(A)	Lock-X(B)	R(A)
Lock-X(B)	Lock-X(B)	Lock-X(B)		Unlock(A)
R(B)	Unlock(A)	R(B)		R(B)
Unlock(A)	R(B)	W(B)		Lock-X(B)
W(B)	W(B)	Commit		R(B)
Unlock(B)	Commit	Unlock(A)		W(B)
Commit	Unlock(B)	Unlock(B)		Unlock(B)

Basic 2PL

Strict 2PL

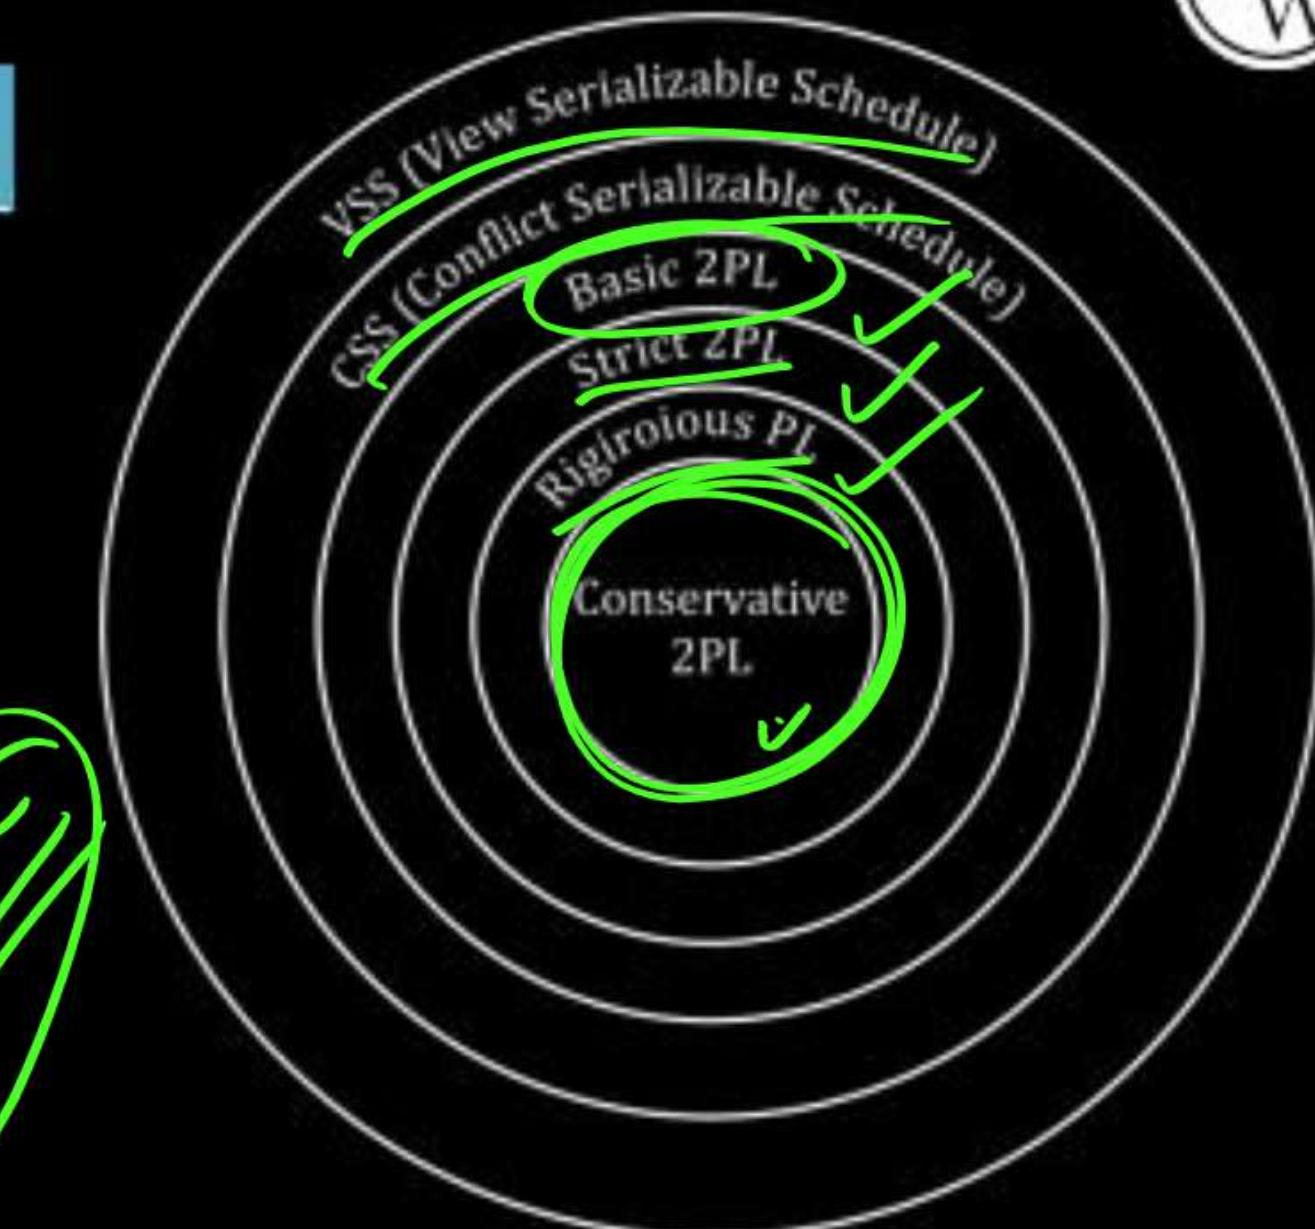
R(A)

+ Strict
+ Basic 2PL

R(B)

W(B)

Not 2PL

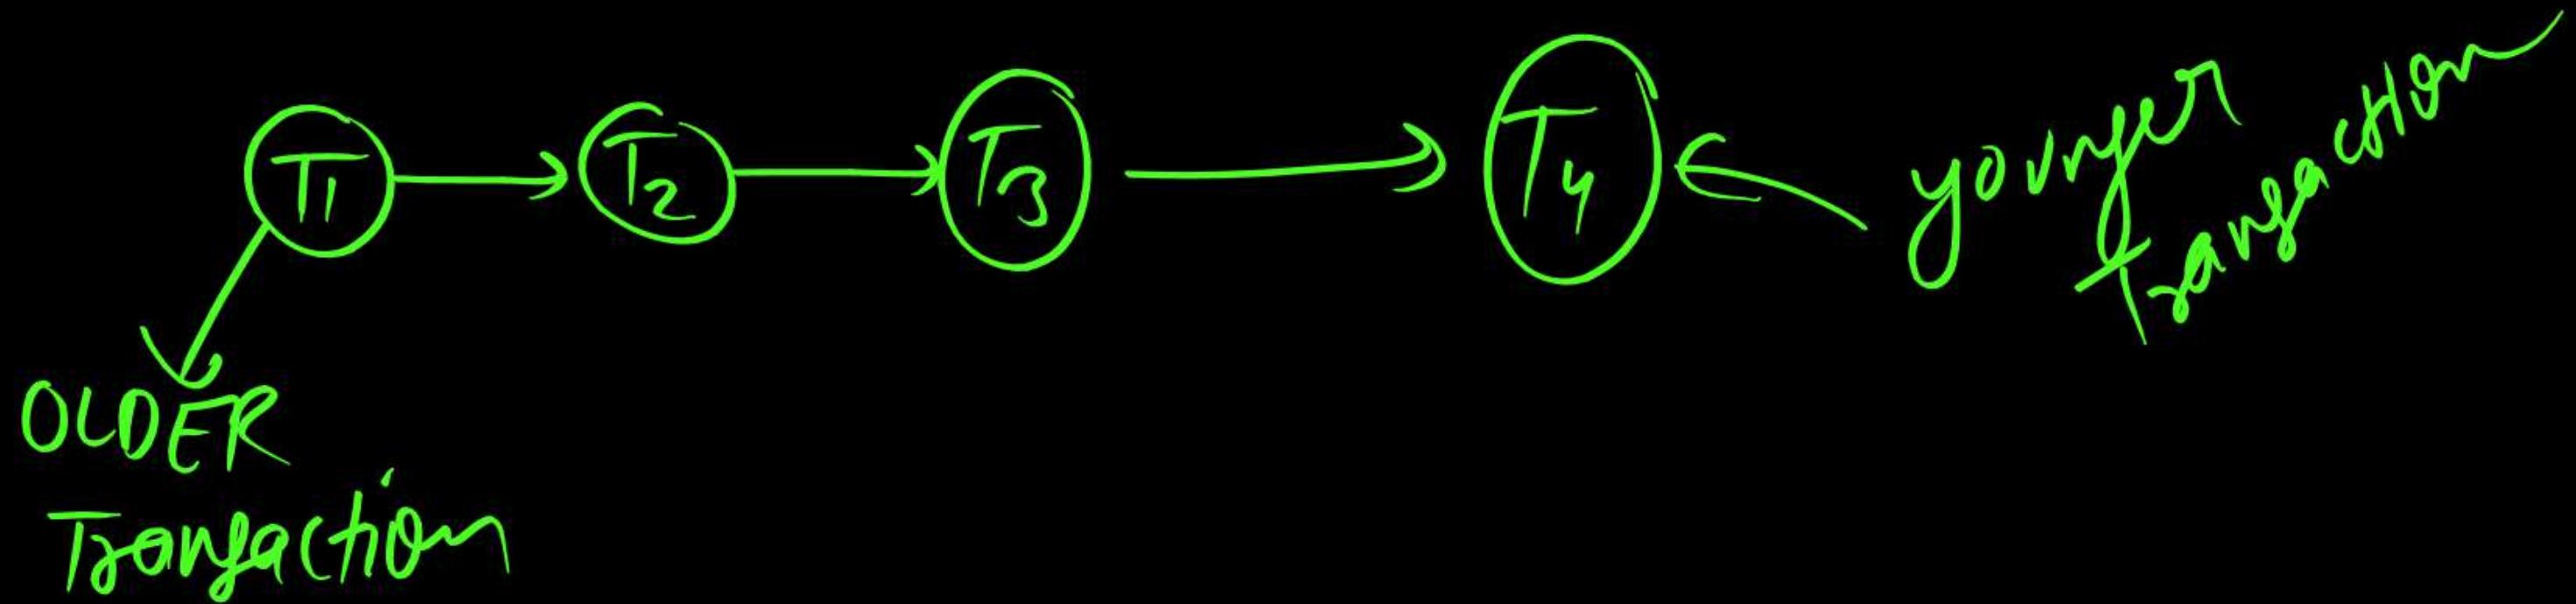




TIMESTAMP BASED CONCURRENCY CONTROL

Time Stamp
Unique Value

(10) (20) (30) (40)
 T_1 T_2 T_3 T_4



$$\begin{aligned} TS(T_1) &= 10 \\ TS(T_2) &= 20 \\ TS(T_3) &= 30 \\ TS(T_4) &= 40 \end{aligned}$$

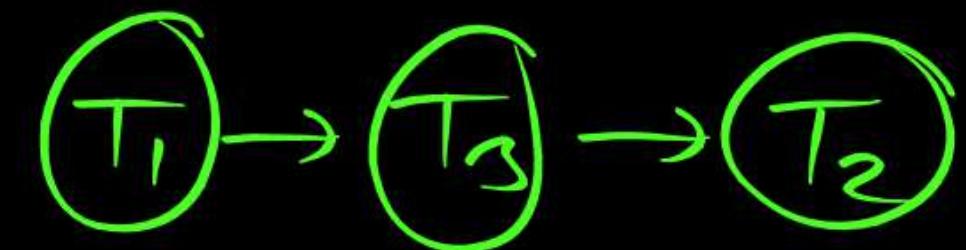
Timestamp-Based Protocols

$TS(T_i)$

- Each transaction T_i is issued a timestamp $TS(T_i)$ when it enters the system.
 - ❖ Each transaction has a unique timestamp
 - ❖ Newer transaction have timestamp strictly greater than earlier ones
 - ❖ Timestamp could be based on a logical counter
 - Real time may not be unique
 - Can use (wall-clock time, logical counter) to ensure
- Timestamp-based protocols manage concurrent execution such that **time-stamp order = serializability order**
- Several alternative protocols based on timestamps

10 30 20
 T_1 $\overline{T_2}$ $\overline{T_3}$

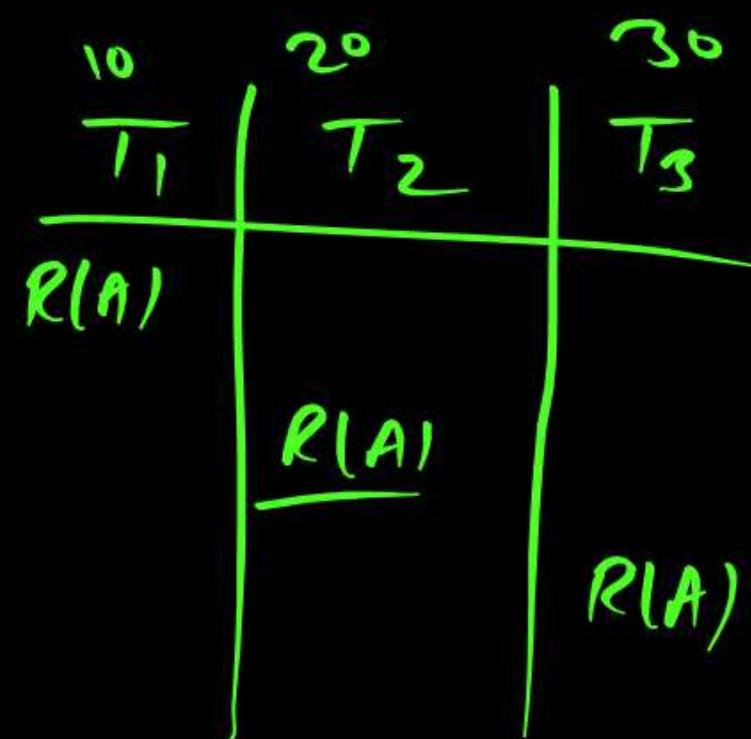
Serializability Order



2 Type of Time Stamp

- ① Read Time Stamp(A) $\Rightarrow \underline{\underline{RTS(A)}}$
- ② Write Time Stamp(A) $\Rightarrow WTS(A)$

RTS(A) : max Time Stamps [R(A)]
successfully



$$RTS(A) = 10$$

$$RTS(A) = 20 \quad [\max(10, 20)]$$

$$RTS(A) = 30 \quad [\max(20, 30)]$$



$$WTS(A) =$$

Timestamp-Based Protocols

The timestamp ordering (TSQ) protocol

- Maintains for each data Q two timestamp values:
 - ❖ W-timestamp(Q) is the largest time-stamp of any transaction that executed write (Q) successfully.
 - ❖ R-timestamp (Q) is the largest time-stamp of any transaction that executed read (Q) successfully.
- Imposes rules on read and write operations to ensure that
 - ❖ any conflicting operations are executed in timestamp order
 - ❖ out of order operations cause transaction rollback

Time Stamp Protocol



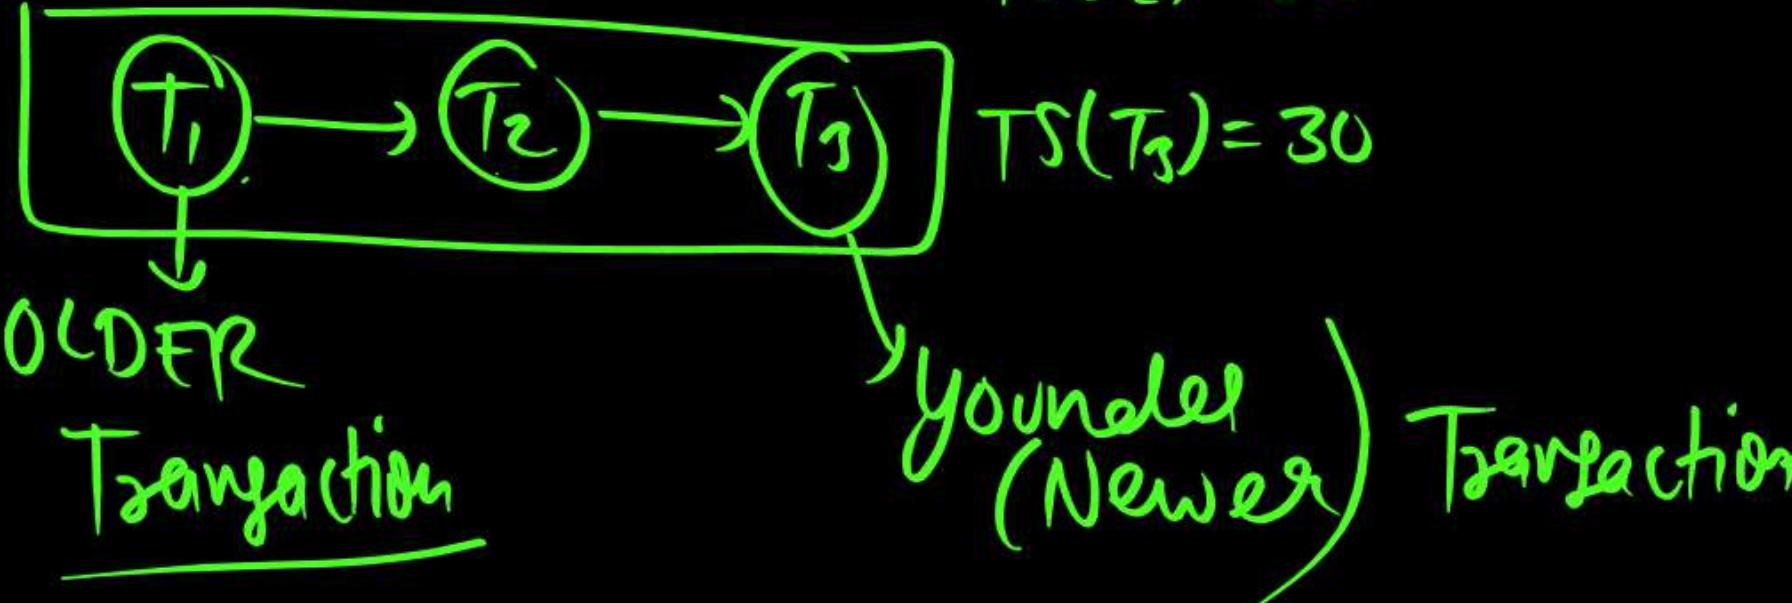
(Fixed)

$$\frac{10}{T_1} \quad \frac{20}{T_2} \quad \frac{30}{T_3}$$

$$TS(T_1) = 10$$

$$TS(T_2) = 20$$

$$TS(T_3) = 30$$



II Time Stamp of Data Item

$$\frac{RTS(Q)}{\omega TS(Q)}$$

max. Transaction stamp
that performs successfully
operation successfully

~~No Transaction Order~~

Same

All Conflict Operation Order

P
W

R-W
W-R
W-W

R-W
W-R
W-W

I T_i : Read(q)

$TS(T_i) < WTS(q)$: Reject, & Rollback

II T_i : Write(q)

$TS(T_i) < RTS(q)$
 $TS(T_i) < WTS(q)$

Write operation & T_i Rollback
Reject

I: T_i - Read(Q) (Transaction T_i Issue R(Q) Operation)

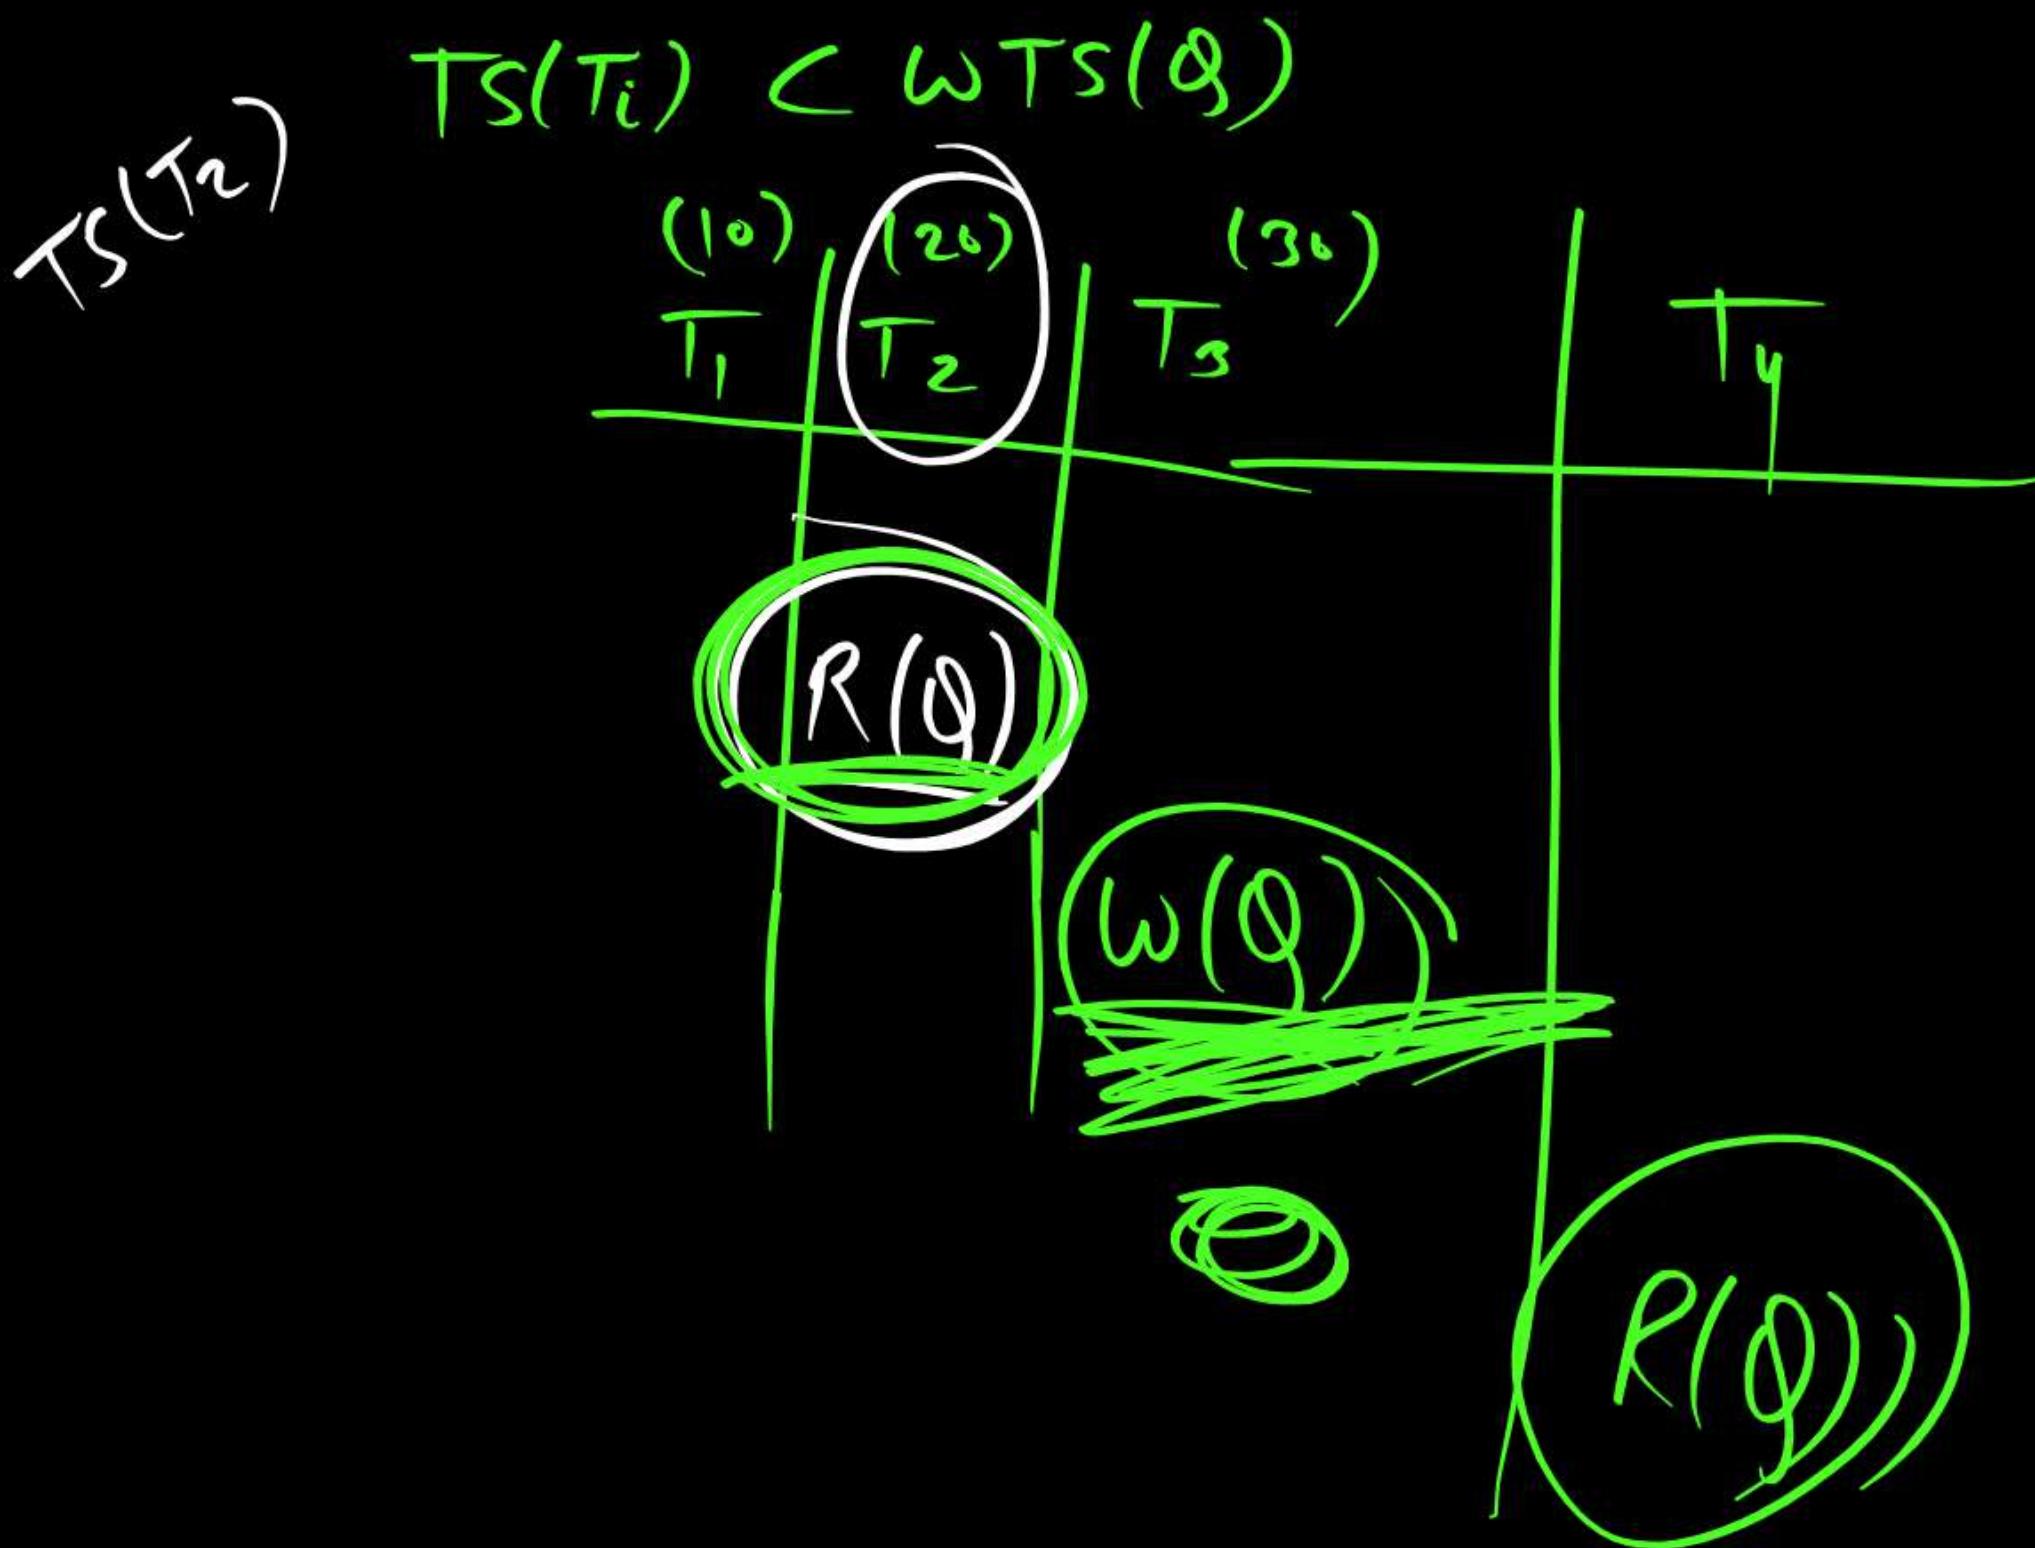
- (i) If TS(T_i) < WTS(Q): Read operation Reject & T_i Rollback.
- (ii) If TS(T_i) ≥ WTS(Q): Read operation is allowed
and Set Read - TS(Q) = max[RTS(Q), TS(T_i)]



II: T_i - Write(Q) (Transaction T_i Issue Write(Q) Operation)

- (i) If TS(T_i) < RTS(Q): Write operation Reject & T_i Rollback.
- (ii) If TS(T_i) < WTS(Q): Write operation Reject & T_i Rollback.
- (iii) Otherwise execute write (Q) operation

Set Read WTS(Q) = TS(T_i)



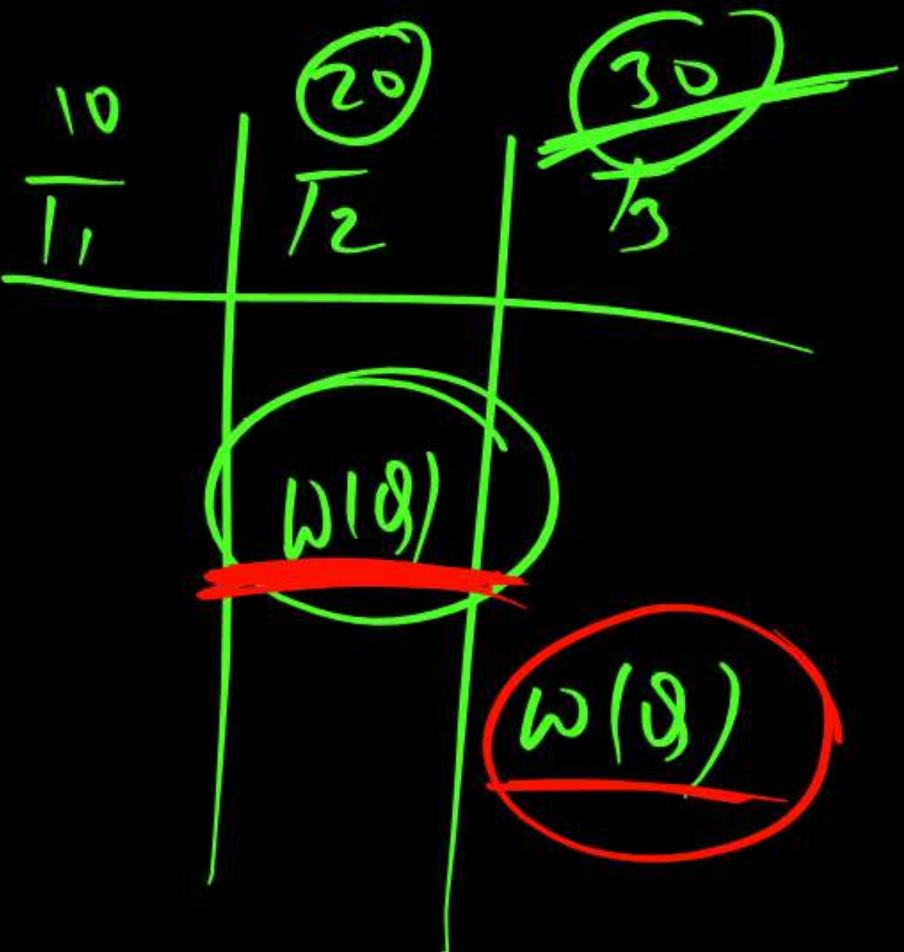
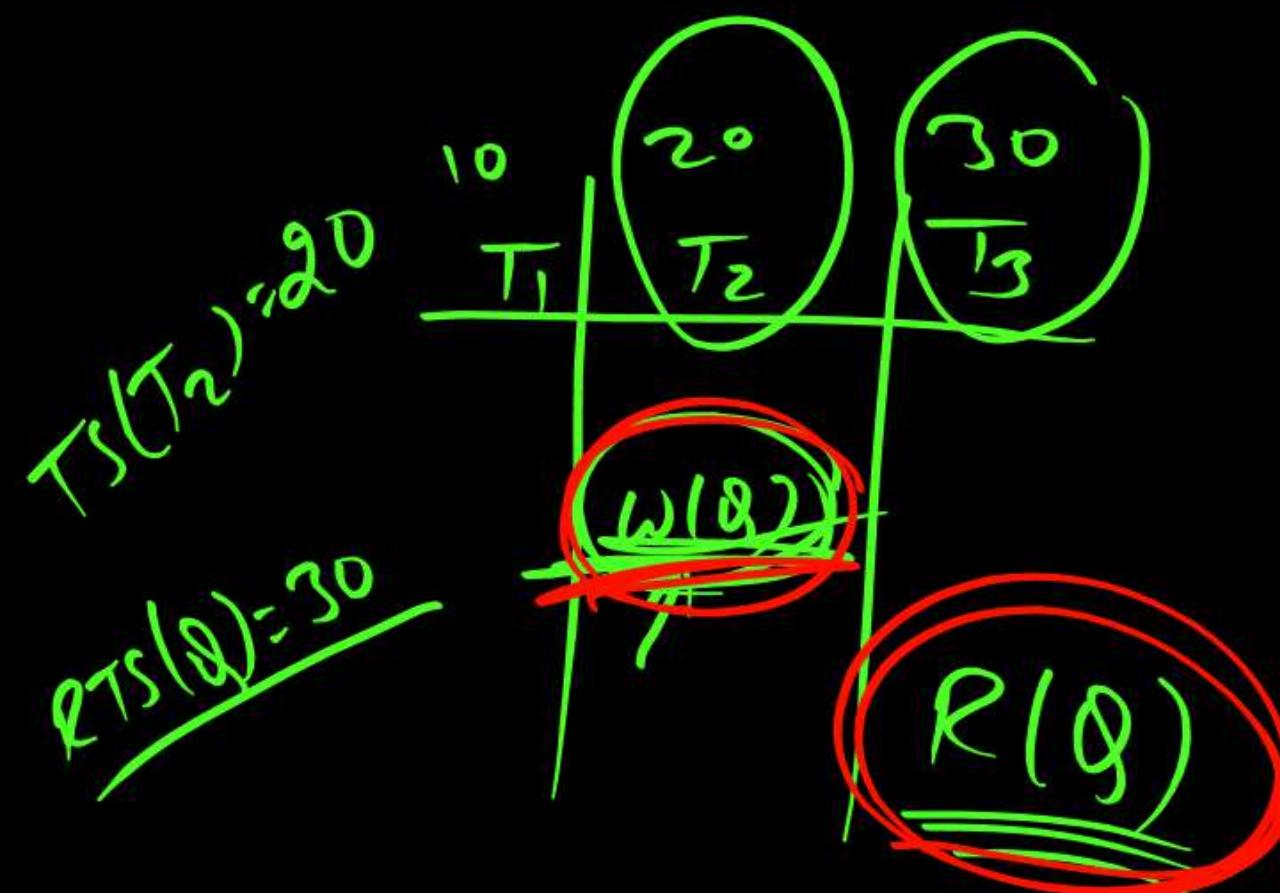
Timestamp-Based Protocols (Cont.)

- Suppose a transaction T_i issues a read (Q)
 1. If $TS(T_i) \leq W\text{-timestamp}(Q)$, then T_i needs to read a value of Q that was already overwritten.
 - Hence, the read operation is rejected, and T_i is rolled back.
 2. If $TS(T_i) \geq W\text{-timestamp}(Q)$, then the read operation is executed, and $R\text{-timestamp}(Q)$ is set to.
 $\max(R\text{-timestamp}(Q), TS(T_i))$.

Timestamp-Based Protocols (Cont.)

- Suppose a transaction T_i issues **write(Q)**
 1. If $TS(T_i) < R\text{-timestamp}(Q)$, then the value of Q that T_i is producing was needed previously, and the system assumed that the value would never be produced.
 - Hence, the **write** operation is rejected, and T_i is rolled back.
 2. If $TS(T_i) < W\text{-timestamp}(Q)$, then T_i is attempting to write an obsolete value of Q.
 - Hence, this **write** operation is rejected, and T_i is rolled back.
 3. Otherwise, the **write** operation is executed, and $W\text{-timestamp}(Q)$ is set to $TS(T_i)$.

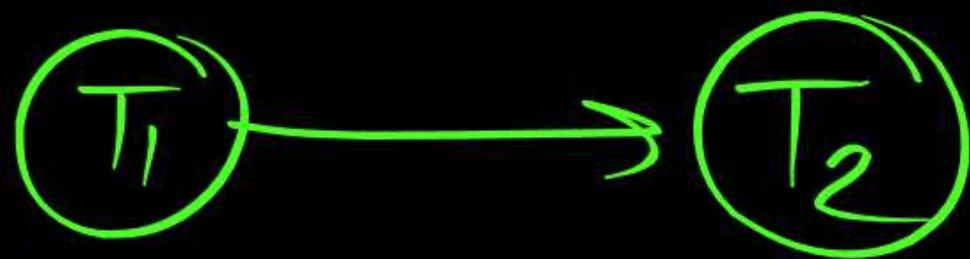
P
W



If

Transaction order

⑧



(T_1 followed by T_2)

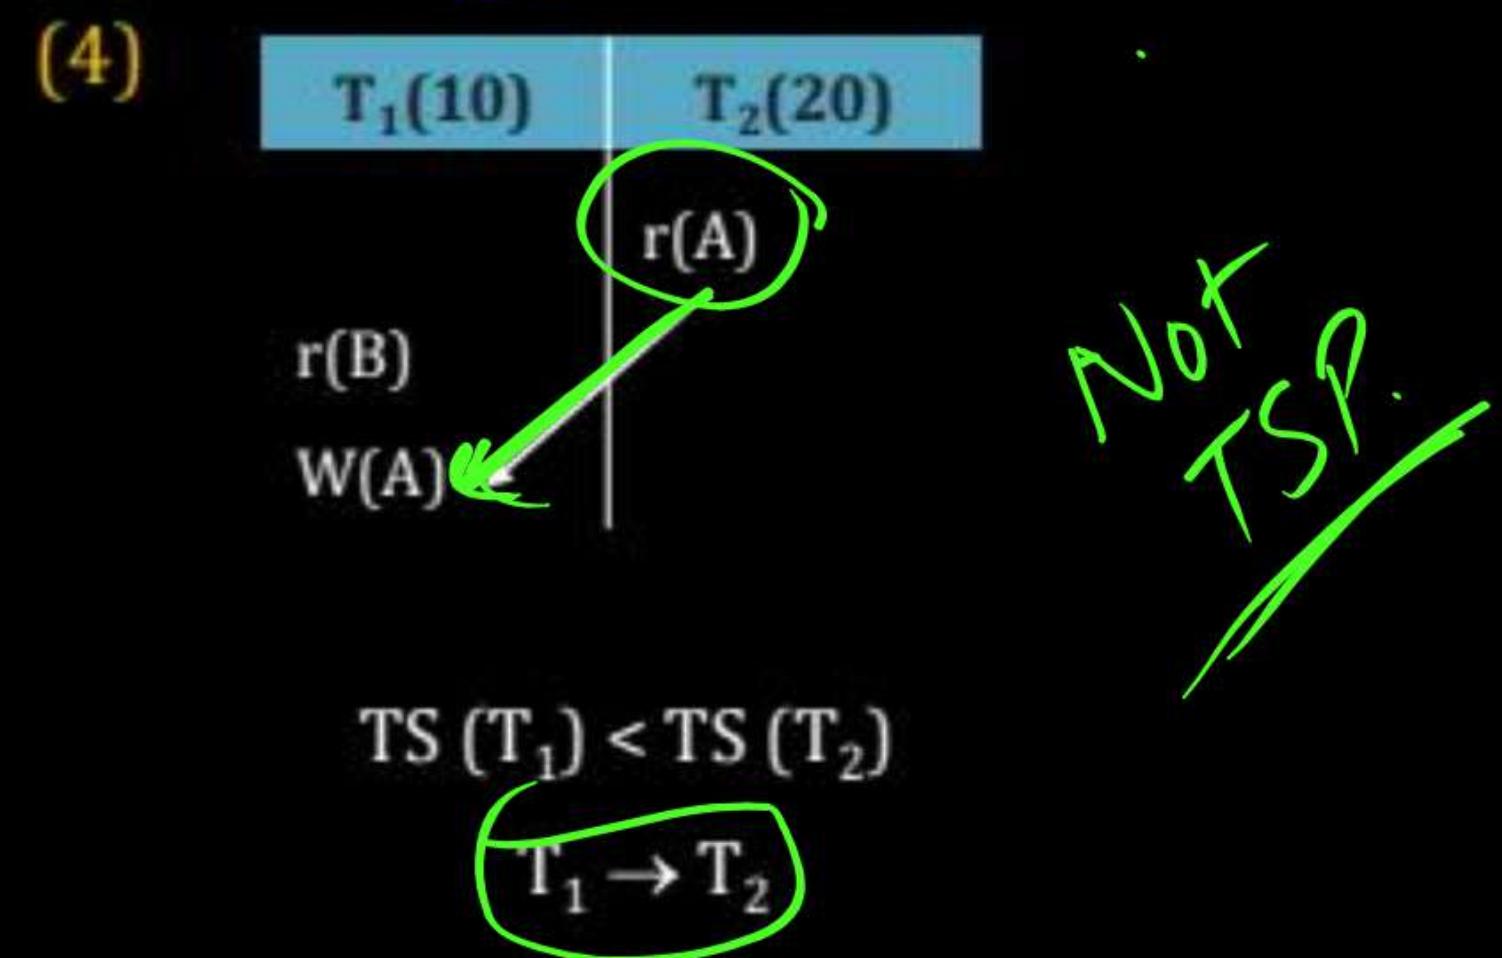
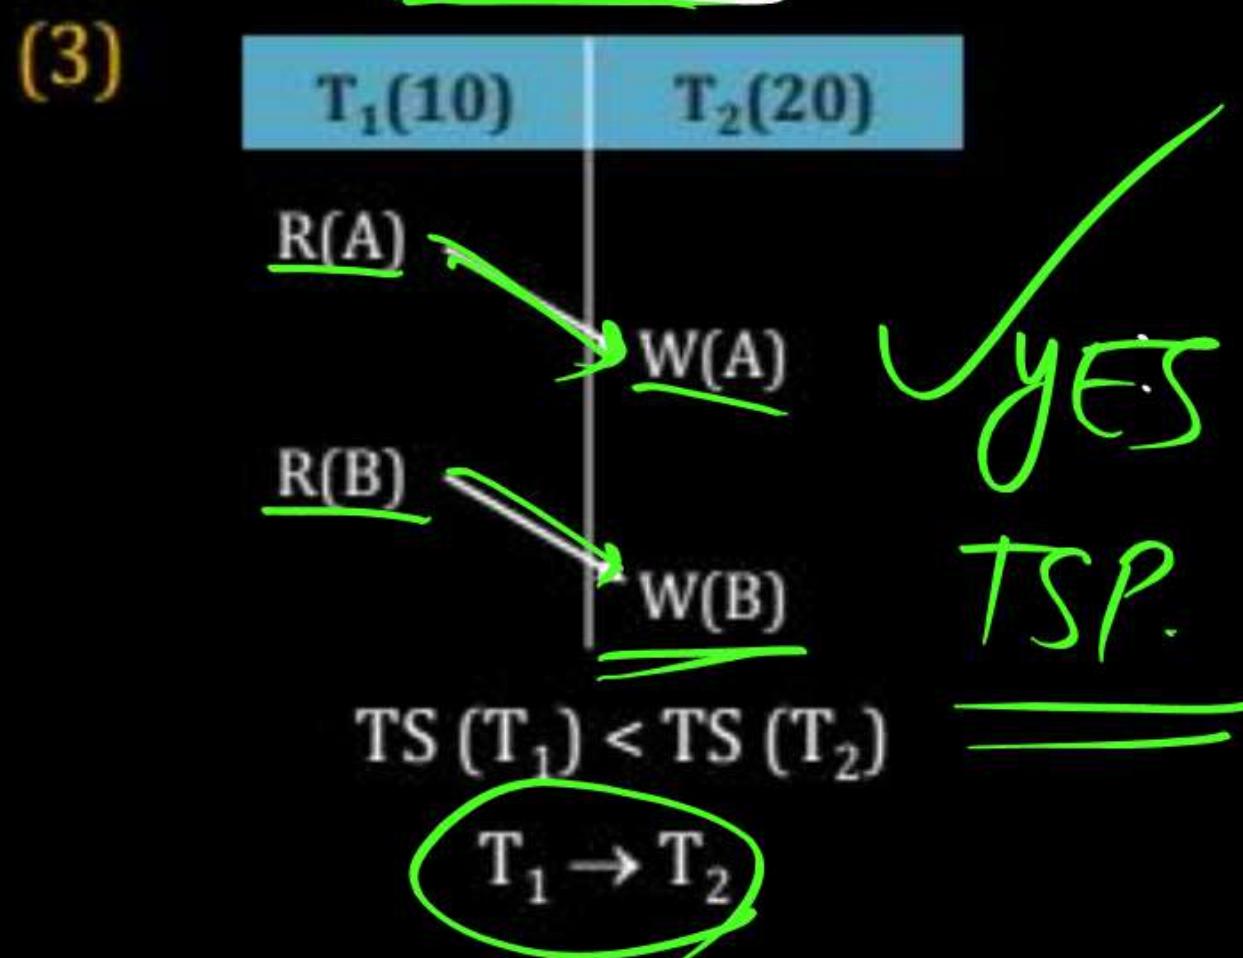
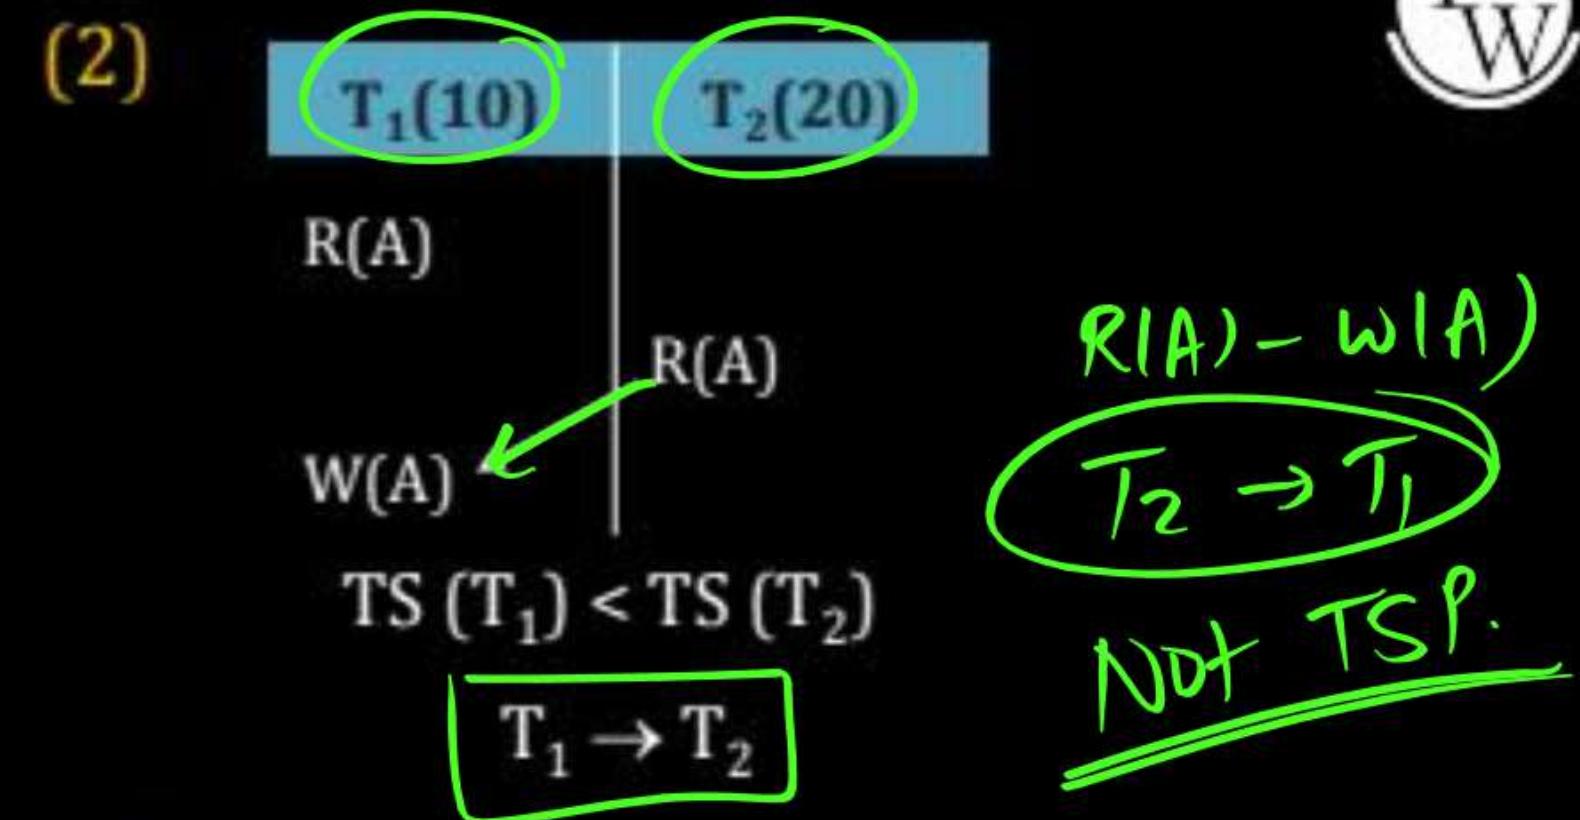
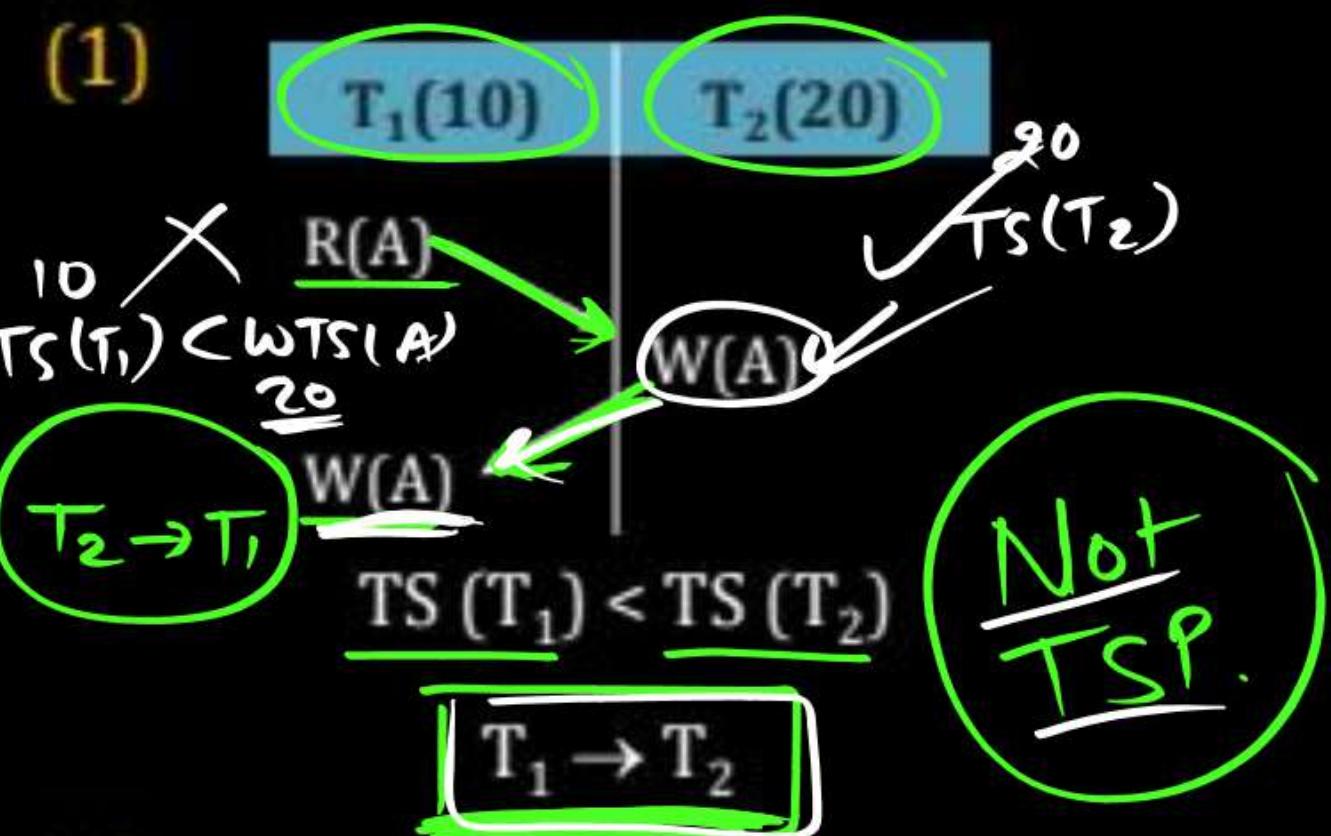
All conflict operation
order them
its allowed by

TSP.

⑨

All conflict operation
order must be $T_1 \rightarrow T_2$

P
W



Thomas Write Rule (View Serializability)

$T_i : \text{Read}(Q)$

$TS(T_i) < WTS(Q)$; Reject & Rollback

Otherwise Read operation allowed.

$T_i : \text{Write}(Q)$

① $TS(T_i) < RTS(Q)$; Reject & Rollback

② $TS(T_i) < WTS(Q)$; Ignored & No Rollback

Thomas' Write Rule

- ❑ Modified version of the timestamp-ordering protocol in which obsolete write operations may be ignored under certain circumstances.
- ❑ When T_i attempts to write data item Q , if $TS(T_i) < W\text{-timestamp}(Q)$, then T_i is attempting to write an obsolete value of $\{Q\}$.
 - ❖ Rather than rolling back T_i as the timestamp ordering protocol would have done, this {write} operation can be ignored.
- ❑ Otherwise this protocol is the same as the timestamp ordering protocol.
- ❑ Thomas' Write Rule allows greater potential concurrency.
 - ❖ Allows some view-serializable schedules that are not conflict-serializable.

Thomas Write Rule (View Serializability)

1. $\text{TS}(T_i) < \text{RTS}(Q)$: Rollback
2. $\text{TS}(T_i) < \text{WTS}(Q)$: Write operation is Ignored and No Roll back

Same as TSP

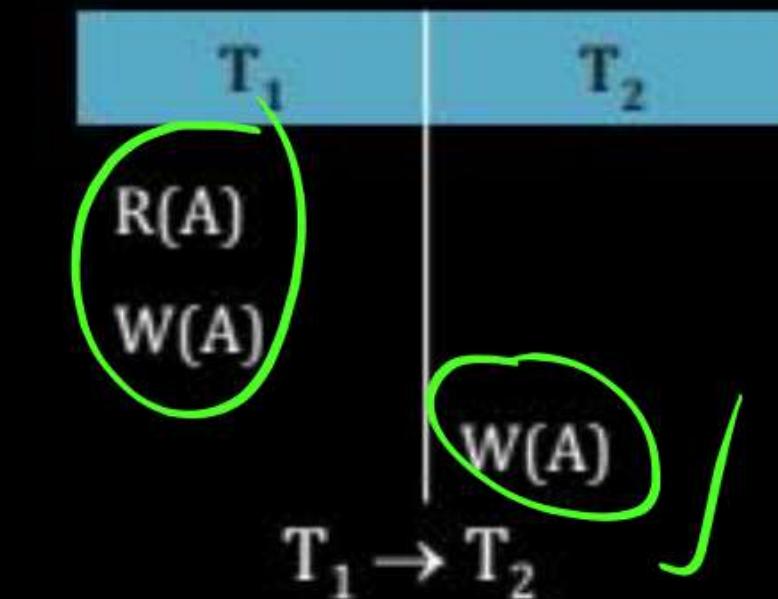
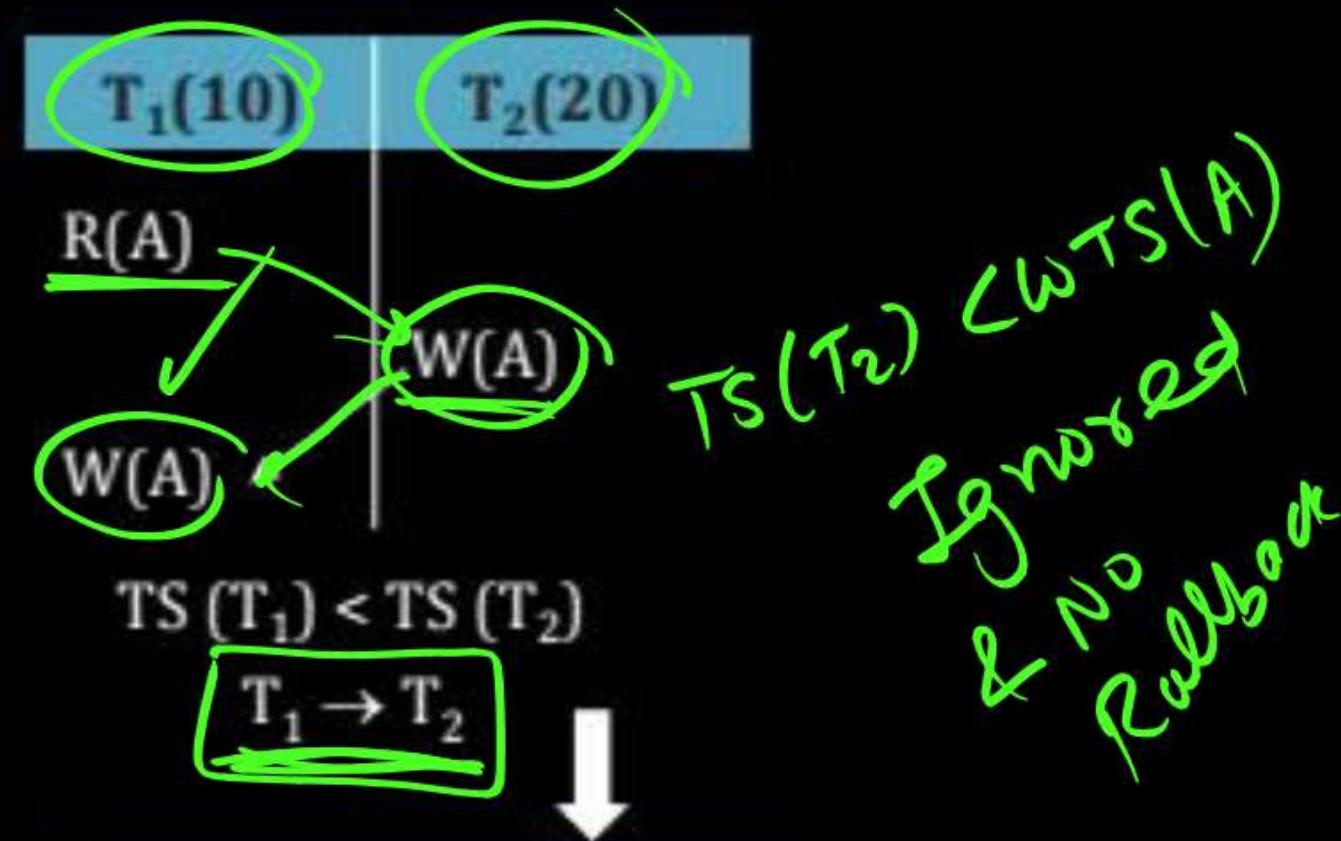
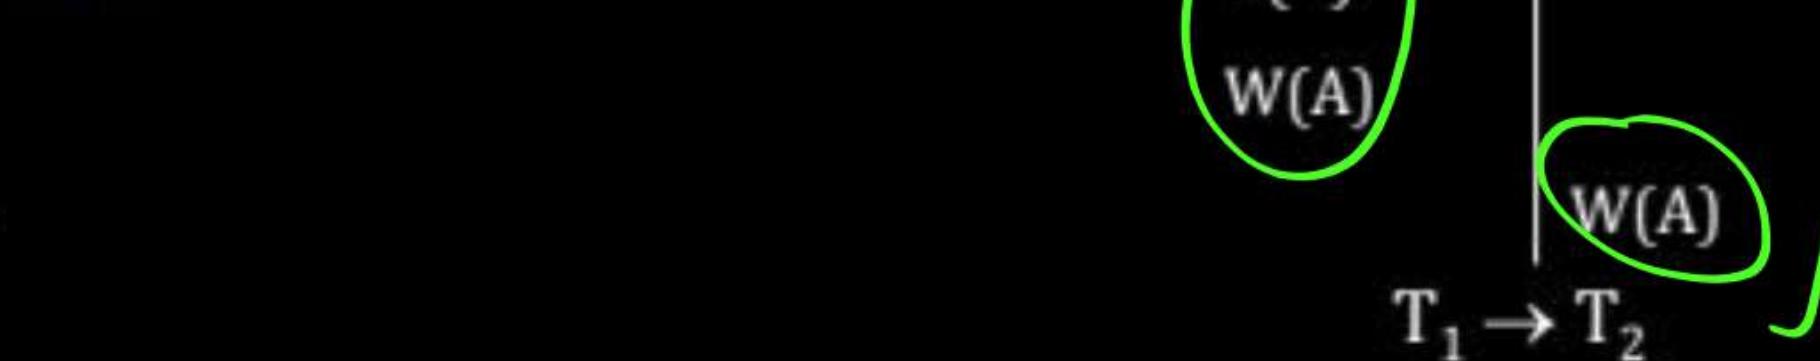
Time Stamp Protocol: Ensure serializability
deadlock free but starvation possible

Deadlock Prevention Algorithm

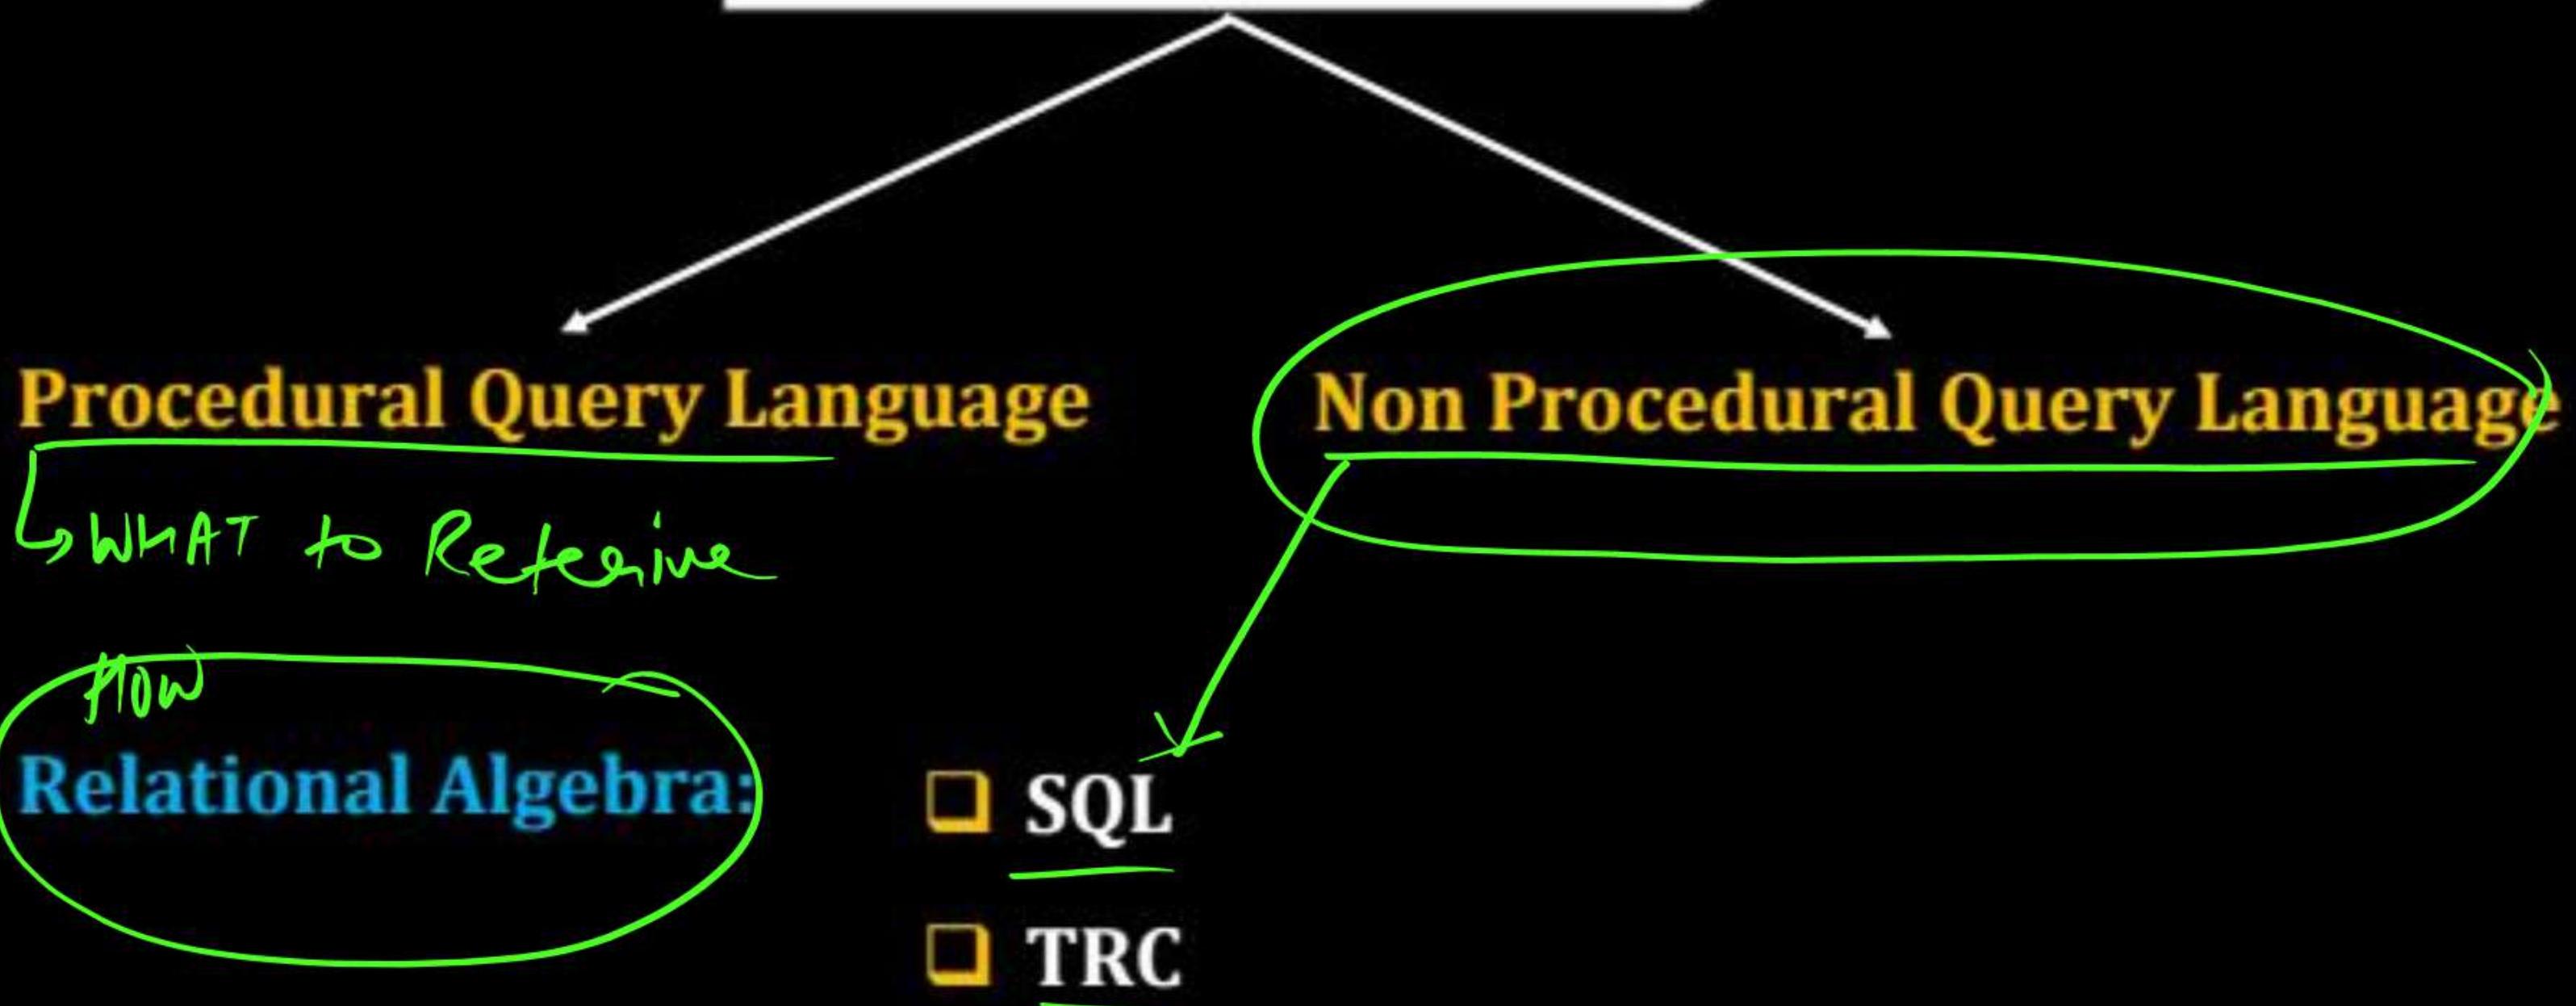
(1) Wait-Die



(2) Wound-wait



Query Language



Relational Algebra

- Procedural language
- Six basic operators
 - ❖ Select: σ
 - ❖ Project: π
 - ❖ Union: \cup
 - ❖ Set difference: -
 - ❖ Cartesian product: \times
 - ❖ Rename: ρ
- The operators take one or two relations as inputs and produce a new relation as a result.

Derived operators

- ① JOIN (\bowtie)
- ② Division [\div]
- ③ Intersection (\cap)

$$R \cap S = R - (R - S)$$

$$R = [1, 2, 3, 4, 5]$$

$$S = [3, 4, 5, 6, 7]$$

$$R - S = [1, 2]$$

$$R \cap S = [3, 4, 5]$$

$$R - (R - S)$$

$$R - (R - S) \Rightarrow [1, 2, 3, 4, 5] - (1, 2) = [3, 4, 5]$$

Note

Basic Idea of Query Language

Executed

↳ Table by Table (Row by Row)

Reserves (R_1)

<u>Sid</u>	<u>Bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

Sailors(S_1)

<u>Sid</u>	<u>Sname</u>	<u>Rating</u>	<u>age</u>
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

Sailors(S_2)

<u>Sid</u>	<u>Sname</u>	<u>Rating</u>	<u>age</u>
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

Selection (σ)

Syntax

$$\sigma_{\text{Condition}}(\text{Relation})$$

If Select tuples Based on specified Condition

(e.g.)

$$\sigma_{CGPA > 8}(\text{STUDENT})$$

Projection (π)

Syntax

$$\pi_{\text{Attribute} \text{ (or)} \text{ Attribute list}}(\text{Relation})$$

If Select Attribute / Attribute List.

$$\pi_{\text{Sid}}(\text{Student})$$
$$\pi_{\text{Sid, Contact}}(\text{STUDENT})$$

Selection (σ)

Q.

 $\sigma_{\text{rating} > 8} (S_2)$

Output:

Sid	Sname	Rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

Sailors (R_2)

Sid	Sname	Rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

Q.

 $\pi_{\text{sname, rating}} (\sigma_{\text{rating} > 8} (S_2))$

Output:



Sname	Rating
yuppy	9
rusty	10

Projection (π)
↓
'Distinct'

Projection (π)

Q.

 $\pi_{age}(S_2)$

Output:

age
35.0
55.5

Sailors (S_2)

Sid	Sname	Rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

Q.

 $\pi_{\underline{\text{sname, rating}}}(S_2)$

Output:

Sname	Rating
yuppy	9
lubber	8
guppy	5
rusty	10

Set operationUnion

- ① Arity Same
- ② Range Similar

Attribute

 \sqsubseteq \sqsupseteq
$$R = \begin{array}{|c|c|c|} \hline 1 & 2 & 3 \\ \hline \end{array} \times \begin{array}{|c|c|} \hline 0 & 1 \\ \hline \end{array}$$

RollNo	Name
X	X

Bronon	Contact
X	X

Union Operation

- Notation: $r \cup s$ [U]
- Defined as :
$$r \cup s = \{t | t \in r \text{ or } t \in s\}$$
- For $r \cup s$ to be valid.
 1. r, s must have the same arity (same number of attributes)
 2. The attribute domains must be compatible (example: 2nd column of r deals with the same type of values as does the 2nd column of s)

Example:

To find all courses taught in the Fall 2009 semester, or in the Spring 2010 semester, or in both.

$$\pi_{\text{course_id}}(\sigma_{\text{semester} = \text{"Fall"} \wedge \text{year} = 2009} (\text{section})) \cup \\ \pi_{\text{course_id}}(\sigma_{\text{semester} = \text{"Spring"} \wedge \text{year} = 2010} (\text{section}))$$

Set Difference Operation

- ❑ Notation: $r - s$
- ❑ Defined as :

$$r - s = \{t \mid t \in r \text{ and } t \notin s\}$$

- ❑ Set differences must be taken between compatible relations.
 - ❖ r and s must have the same arity
 - ❖ attribute domains of r and s must be compatible

Example:

Example: to find all courses taught in the Fall 2009 semester, but not in the Spring 2010 semester

$\pi_{course_id}(\sigma_{semester = "Fall"} \wedge year = 2009 \text{ (section)})$

$\pi_{course_id}(\sigma_{semester = "Spring"} \wedge year = 2010 \text{ (section)})$

Assume Relation R & Relation S consist M & N Tuple Respectively

$R \Rightarrow m$ Table

M N

$S \Rightarrow N$ Table

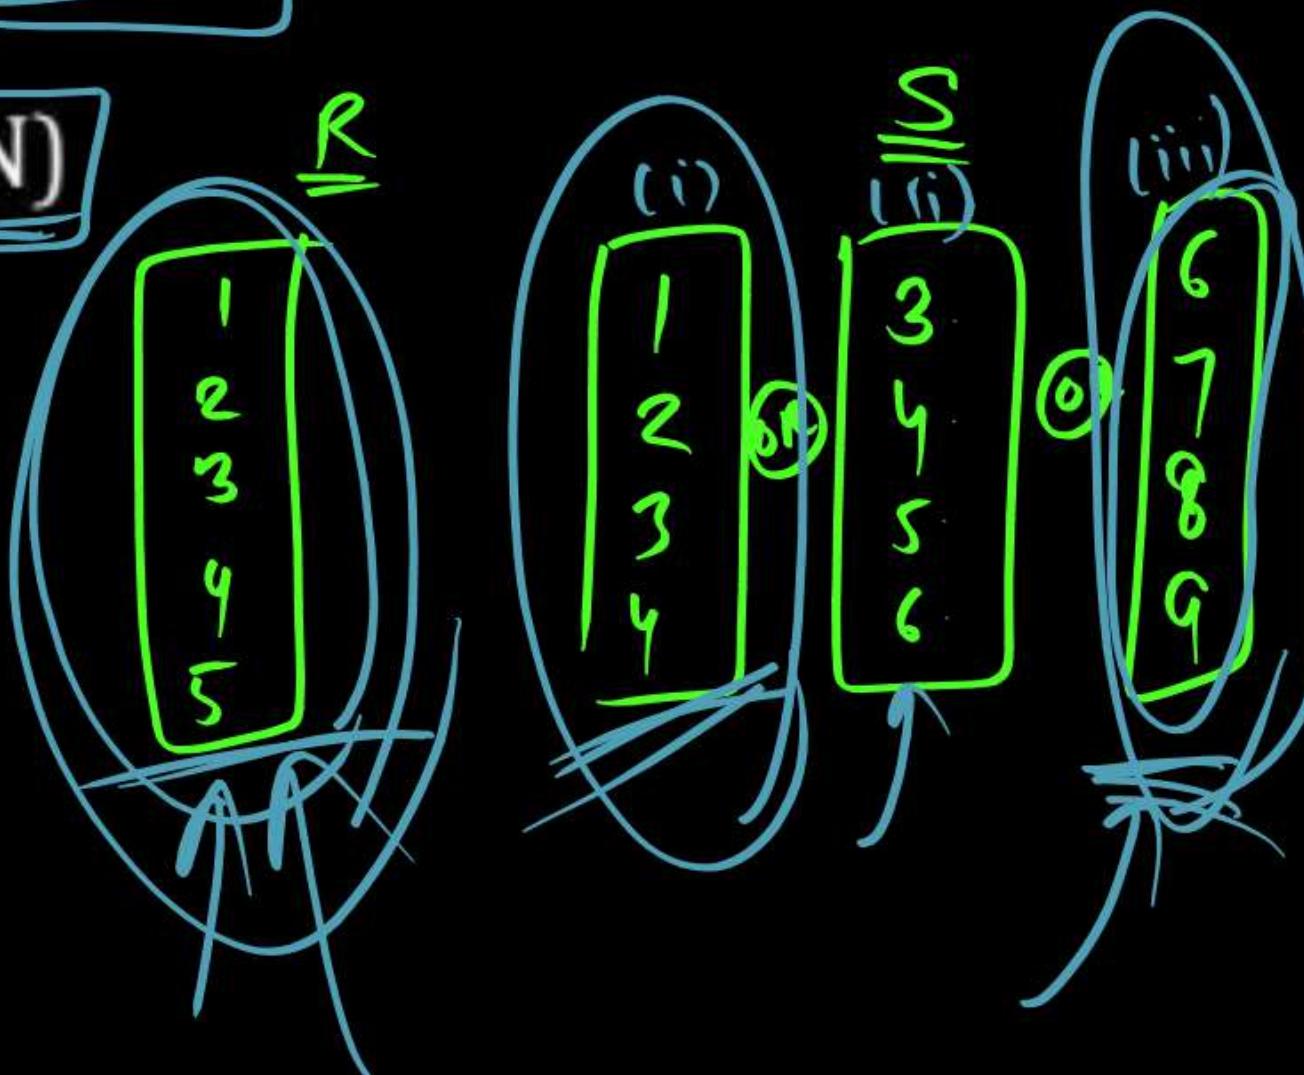
max

(1) Range of tuples in $R \cup S = \boxed{\max(M, N)}$ to $\boxed{M + N}$

(2) Range of tuples in $R \cap S = \phi$ to $\min(M, N)$

(3) Range of tuples in $R - S = \phi$ to M

(4) Range of tuples in $S - R = \phi$ to N



Example:

Sailors (S_1)

S_1

<u>Sid</u>	Sname	Rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

Sailors (S_2)

28 Yuppy
44 Guppy

S_2

<u>Sid</u>	Sname	Rating	age
28	Yuppy	9	35.0
31	Lubber	8	55.5
44	Guppy	5	35.0
58	rusty	10	35.0

5

$S_1 \cup S_2$

2

$S_1 \cap S_2$

$S_1 - S_2 \Rightarrow 1$

1
2
28
44

$S_2 - S_1 \Rightarrow 0$

(1) Union

Sid	Sname	Rating	age	
22	dustin	7	45.0	✓
31	lubber	8	55.5	✓
58	rusty	10	35.0	✓
44	guppy	5	35.0	✓
28	yuppy	9	35.0	✓

 $S1 \cup S2$

(2) Set Difference

Sid	Sname	Rating	age
22	dustin	7	45.0

S1 - S2

(3) Intersection

Sid	Sname	Rating	age
31	lubber	8	55.5
58	rusty	10	35.0

$S1 \cap S2$

CROSS Product $(R \times S)$

R
 n_1 Tuple

C_1 Attribute

S
 n_2 Tuple

C_2 Attribute

$R \times S = n_1 \times n_2$ Tuple
 $C_1 + C_2$ Attribute

Cross - Product

P
W

3 Tuple
4 Attribute

$\Rightarrow 2 \times 3 = 6 \text{ Tuple}$
 $3 + 4 = 7 \text{ Attribute}$

Reserves (R_1)

2 Tuple
3 Attribute

(3)

Sid	bid	day
22	101	10/10/96
58	103	11/12/96

Sailors (S_1)

Sid	Sname	Rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

$R_1 : 2 \text{ Tuple}$
3 Attribute

$S_1 : 3 \text{ Tuple}$
4 Attribute

$$R_1 \times S_1 = 2 \times 3 = 6 \text{ Tuple}$$

$$3 + 4 = 7 \text{ Attribute}$$

$R_1 \times S_1$

(Sid)	Sname	Rating	age	(Sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

Selection (σ)

Projection (π)

① Arity

② Range

Union
Intersection

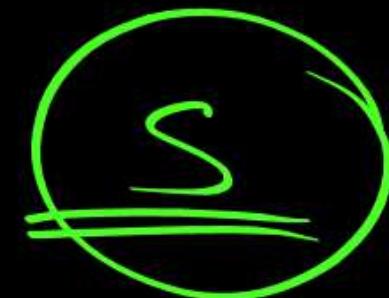
Set Difference

Arity
 R

CROSS Product

$R \times S = n_1 \times n_2$ Tuple
 $C_1 + C_2$ Attribute

CROSS Product ?



$R \times S$ \leftarrow

A circle containing the letters 'S' and 'Y' with three horizontal lines underneath it, all crossed out.

Join Operations

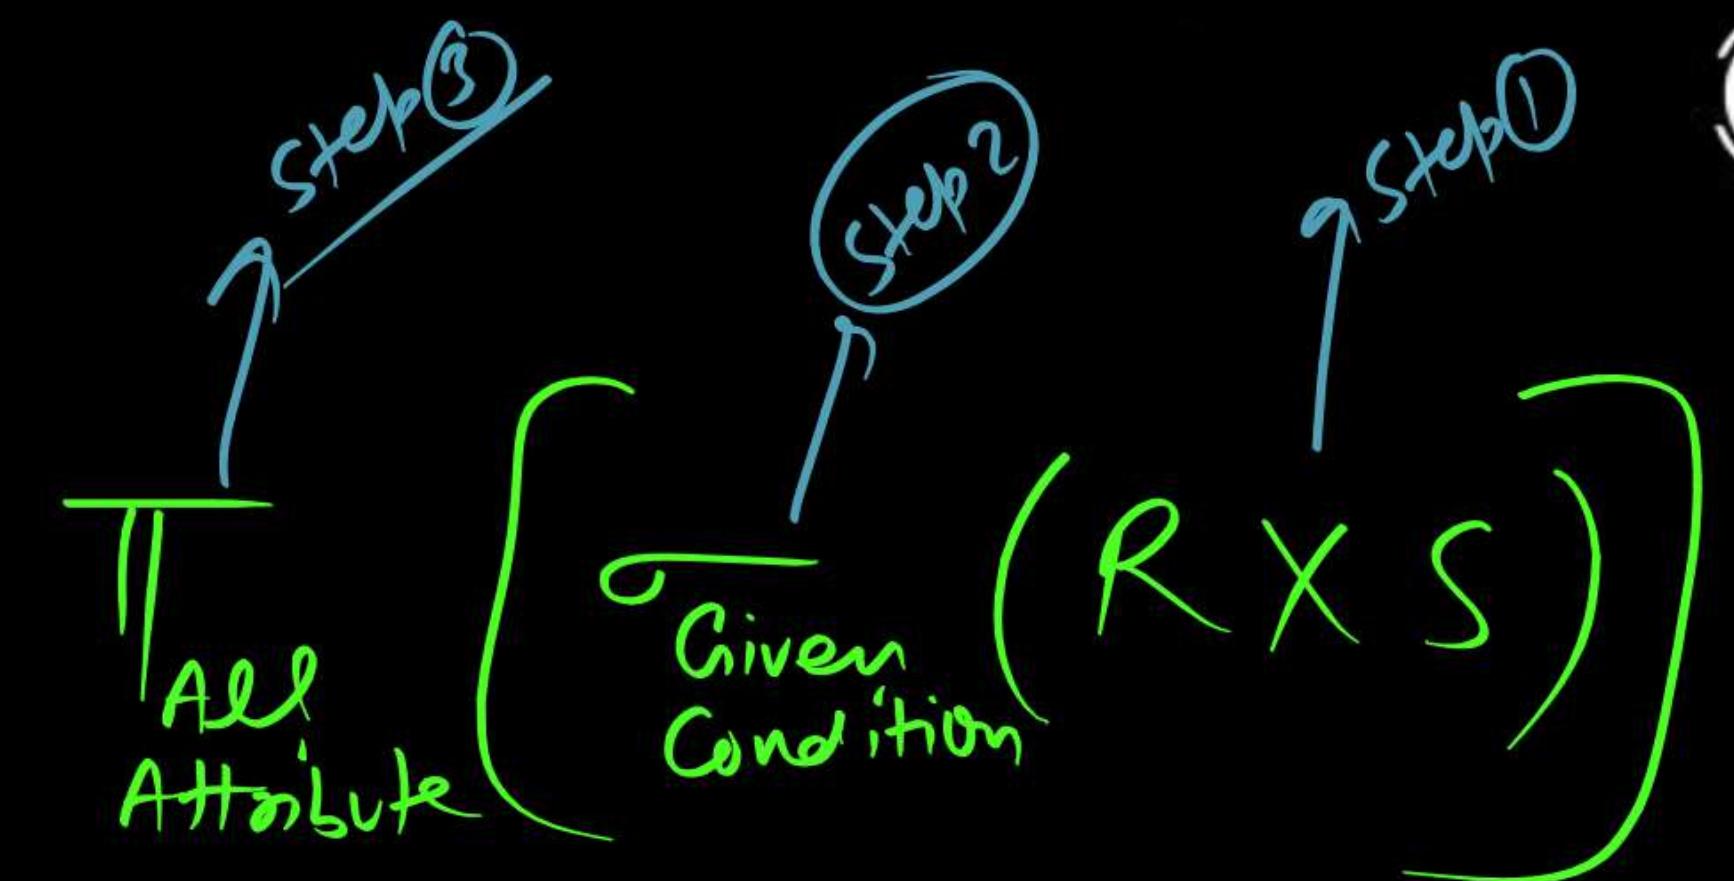
- (1) Conditional Join (\bowtie_c)
- (2) Equi join
- (3) Natural join
- (4) Left outer join
- (5) Right outer join
- (6) Full outer join

JOIN (\bowtie_c) $(R \bowtie_c S)$

① Condition Join

Step 1 CROSS Product of R & S. $(R \times S) \Rightarrow$ ^{n₁ × n₂ Table}
_{Attribute}Step 2 Select the Table which satisfy given conditionStep 3 Projection of Attribute

$R \bowtie_c S =$



Conditional Join

$$R \bowtie_c S = \sigma_C (R \times S)$$

R_1 Sid

(Sid)	Sname	Rating	age	(Sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

$R_1 \times S_1$

S_1 Sid

Conditional Join

$$S1 \bowtie_{S1.sid < R1.sid} R1$$

output

(Sid)	Sname	Rating	age	(Sid)	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	58	103	11/12/96

$$S1 \bowtie_{S1.sid < R1.sid} R1$$

- ❖ Result schema same as that of cross-product.
- ❖ Fewer tuples than cross - product, might be able to compute more efficiently.
- ❖ Sometimes called a theta -join.

Equi – Join

A special case of condition join where the condition c contains only equalities.

$$R_1.Sid = S_1.Sid$$

$R_1 \times S_1$	R_1	S_1	
	R_1	S_1	
	22	dustin	22
	22	dustin	58
	31	lubber	22
	31	lubber	58
	58	rusty	22
	58	rusty	58

Equi Join \Rightarrow Equality Condition

Equi – Join

$S1 \bowtie_{\text{sid}} R1$

Sid	Sname	Rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	11/12/96

- ❖ Result schema similar to cross - product, but only one copy of fields for which equality is specified.
- ❖ Natural join: Equijoin on all common fields.

$R(A B C)$ $S(B C D)$ $R \bowtie S =$

$\Pi_{\text{Distinct Attribute}}$ [equality condition on all common attributes]

 $R \bowtie S =$

$\Pi_{ABCD} [R.B = S.B \wedge R.C = S.C \wedge (R \times S)]$

Natural Join (\bowtie) ($R \bowtie S$)

Step ① CROSS Product of $R \& S$ ($R \times S$)

Step ② Select the tuple which satisfy equality condition
on All Common Attribute

Step 3 Projection of Distinct Attribute

JOIN

① Conditional Join

② Equi Join

③ Natural Join.

④ Left Outer Join

⑤ Right Outer Join

⑥ Full Outer Join.

**THANK
YOU!**

